

Cloud computing and distributed systems

Transactions and concurrency control

Zeynep Yücel

Ca' Foscari University of Venice
zeynep.yucel@unive.it
yucelzeynep.github.io

Introduction

- Transaction Definition: set of operations treated as as an indivisible unit.
 - ▶ Either complete a transaction or fully erase it.
- Goal of transactions: Keep a consistent state of objects during
 - ▶ Multiple transactions
 - ▶ Server crash
- Recoverable Objects: are those that can be recovered after the server crashes.
 - ▶ Stored in volatile or persistent memory.

Introduction

Operations of the *Account* interface

deposit(amount)

deposit *amount* in the account

withdraw(amount)

withdraw *amount* from the account

getBalance() → *amount*

return the balance of the account

setBalance(amount)

set the balance of the account to *amount*

Operations of the *Branch* interface

create(name) → *account*

create a new account with a given name

lookup(name) → *account*

return a reference to the account with the given name

branchTotal() → *amount*

return the total of all the balances at the branch

- Each account is a remote object whose interface, *Account*, provides deposit, withdraw and get/setBalance operations.
- Each branch is a remote object whose interface, *Branch*, provides create, lookup, branchTotal operations.

Introduction

Simple synchronization (without transactions) Atomic operations at the server

- Client operations can interfere.
- Synchronization Without Transactions
 - ▶ Object methods should be designed accordingly in a multi-threaded context.
 - ▶ If they are not designed in multi-threaded context, strange interleavings can happen.

Introduction

Simple synchronization (without transactions) Atomic operations at the server

- Java's "synchronized" keyword controls threads' access to an object (one thread at a time)
 - ▶ When a thread calls a synchronized method, the object is locked, blocking other threads until lock release.
 - ▶ Without synchronization, if concurrent *deposit* operations read a balance, then the incremented value will reflect a single deposit.
- Atomic Operations - No thread interference.
 - ▶ Built using synchronized methods

Introduction

Enhancing client cooperation by synchronization of server operations

- Clients use a server for resource sharing.
 - ▶ some clients update server objects, others access them.
- Synchronized access enables safe resource sharing.
- But communication between threads may be necessary for completion.
- Achieved with *wait* and *notify*.
 - ▶ Wait Method - Suspends own execution to allow others to execute.
 - ▶ Notify Method - Alerts waiting threads.
- Atomic access even during waits.

Introduction

Enhancing client cooperation by synchronization of server operations

- Consider shared object *Queue*
 - ▶ *first* method - Removes first object.
 - ▶ *append* method - Adds to queue.
 - ▶ Call *first* on empty queue - Wait for addition.
 - ▶ *append* calls *notify* - Alerts waiting threads.
 - ▶ If threads sync with wait and notify, server answers requests that cannot immediately be satisfied.
- Without synchronization, client need to poll (inefficient and unfair).

Introduction

Failure model for transactions

- Claims of the model:
 - ▶ Predictable faults can be handled.
 - ▶ Errors can be detected and managed before issues arise.
- Permanent Storage Failures:
 - ▶ Writes may fail - Incorrect or missing data.
 - ▶ Reads can detect bad data (e.g. by checksums).
- Server Crashes:
 - ▶ New process replaces it (no recovery from any volatile memory).
 - ▶ Recovery uses permanent storage and info from other processes.
 - ▶ Crashes happen also during recovery
- Communication Issues:
 - ▶ Messages may suffer delays, duplication, loss, or corruption.
 - ▶ Checksums can detect corruption.

Transactions

Transaction T:

a.withdraw(100);

b.deposit(100);

c.withdraw(200);

b.deposit(200);

- Client operations on accounts A, B, C
- Transfer via withdrawal followed by deposit

- \$100 transferred from A to B
- \$200 transferred from C to B

Transactions

All or nothing, isolation

- Transactions apply to recoverable objects and are atomic.
- Atomicity: All or Nothing, i.e. complete or no effect
 - ▶ Failure Atomicity: Effects are atomic
 - ▶ Durability: successful transactions are stored in permanent memory
- Isolation: No interference between parallel transactions, intermediate effects hidden

Transactions

- To satisfy all-or-nothing and isolation requirements,
 - ▶ Objects must be recoverable
 - ▶ Changes must be stored in permanent storage
 - ▶ Operations must be synchronized
- Serial execution of transactions
 - ▶ Unacceptable in most cases

Transactions

Acid properties

- ACID properties mnemonic
 - ▶ Atomicity: All or nothing
 - ▶ Consistency: Transition between two consistent states
 - ▶ Isolation: No mutual effects between transactions
 - ▶ Durability: Persist post-commit

Transactions

Operations in the *Coordinator* interface

openTransaction() → *trans*;

Starts a new transaction and delivers a unique TID *trans*. This identifier will be used in the other operations in the transaction.

closeTransaction(trans) → (*commit*, *abort*);

Ends a transaction: a *commit* return value indicates that the transaction has committed; an *abort* return value indicates that it has aborted.

abortTransaction(trans);

Aborts the transaction.

- Recoverable object server with transaction capabilities
- Coordinator manages transactions
 - ▶ Start transaction with *openTransaction* and receive Transaction Identifier (TID)
 - ▶ End transaction with *closeTransaction*
 - ▶ Abort with *abortTransaction*
- Transactions need cooperation among client, objects, coordinator
- Clients send TID with each operation on recoverable objects.

Transactions

<i>Successful</i>	<i>Aborted by client</i>	<i>Aborted by server</i>
<i>openTransaction</i>	<i>openTransaction</i>	<i>openTransaction</i>
<i>operation</i>	<i>operation</i>	<i>operation</i>
<i>operation</i>	<i>operation</i>	<i>operation</i>
•	•	server aborts
•	•	<i>transaction</i> →
<i>operation</i>	<i>operation</i>	<i>operation ERROR</i> <i>reported to client</i>
<i>closeTransaction</i>	<i>abortTransaction</i>	

- A transaction is either successful or aborted.
 - ▶ Successful: *closeTransaction* by client, changes become permanent, reply is *committed*.
 - ▶ Aborted: no effects visible.
- Two ways to abort:
 - ▶ Client aborts the transaction calling *abortTransaction*.
 - ▶ Server aborts the transaction.

Transactions

Client actions related to server process crashes

- Server actions against crashes
 - ▶ if server crashes, replace server process, abort uncommitted transactions, restore objects.
 - ▶ if client crashes, give expiry times to transactions.
- Client actions against crashes
 - ▶ if server crashes, operation returns an exception after a timeout.
 - ▶ Even if the server process is replaced, the old transaction is not valid.

Transactions

Concurrency control - The lost update problem

- Lost Update Problem.
- Concurrent updates lead to loss of information.
- Bank accounts A (\$100), B (\$200), C (\$300).
 - ▶ Transaction T: A to B.
 - ▶ Transaction U: C to B.
 - ▶ Amount is 10% of B.
- Expected balance: \$242 ($200 \rightarrow 220 \rightarrow 242$).
- With concurrent execution, balance becomes \$220.

Transactions

Inconsistent retrievals

Transaction T:	Transaction U:
<i>balance = b.getBalance();</i> <i>b.setBalance(balance*1.1);</i> <i>a.withdraw(balance/10)</i>	<i>balance = b.getBalance();</i> <i>b.setBalance(balance*1.1);</i> <i>c.withdraw(balance/10)</i>
<i>balance = b.getBalance();</i> \$200	<i>balance = b.getBalance();</i> \$200
	<i>b.setBalance(balance*1.1);</i> \$220
<i>b.setBalance(balance*1.1);</i> \$220	
<i>a.withdraw(balance/10)</i> \$80	<i>c.withdraw(balance/10)</i> \$280

- Balance of B: \$200
- T and U increase B's balance by \$20 and \$22
- U's update lost

Transactions

Inconsistent retrievals

Transaction V:		Transaction W:	
<i>a.withdraw(100)</i>		<i>aBranch.branchTotal()</i>	
<i>b.deposit(100)</i>			
<i>a.withdraw(100);</i>	\$100	<i>total = a.getBalance()</i>	\$100
		<i>total = total + b.getBalance()</i>	\$300
		<i>total = total + c.getBalance()</i>	
<i>b.deposit(100)</i>	\$300	•	
		•	

- Inconsistent Retrievals Problem:
 - ▶ Transaction V: transfers from A to B \$100
 - ▶ Transaction W: sum of balances of all account in the branch
 - ▶ Accounts A and B: \$200
 - ▶ BranchTotal: \$300, wrong

Transactions

Inconsistent retrievals - Serial Equivalence

Transaction T:		Transaction U:	
<i>balance = b.getBalance()</i>		<i>balance = b.getBalance()</i>	
<i>b.setBalance(balance*1.1)</i>		<i>b.setBalance(balance*1.1)</i>	
<i>a.withdraw(balance/10)</i>		<i>c.withdraw(balance/10)</i>	
<i>balance = b.getBalance()</i>	\$200	<i>balance = b.getBalance()</i>	\$220
<i>b.setBalance(balance*1.1)</i>	\$220	<i>b.setBalance(balance*1.1)</i>	\$242
<i>a.withdraw(balance/10)</i>	\$80	<i>c.withdraw(balance/10)</i>	\$278

- Serial Equivalence:
 - ▶ If two transactions are correct individually, their sequential execution will also be correct.
 - ▶ Interleaving operations can be serially equivalent if the overall effect remains the same.
- Avoid lost updates
- Later transaction reads updated

Transactions

Inconsistent retrievals - Serial Equivalence

Transaction V:		Transaction W:	
<i>a.withdraw(100);</i> <i>b.deposit(100)</i>		<i>aBranch.branchTotal()</i>	
<i>a.withdraw(100);</i>	\$100	<i>total = a.getBalance()</i>	\$100
<i>b.deposit(100)</i>	\$300	<i>total = total + b.getBalance()</i>	\$400
		<i>total = total + c.getBalance()</i>	
		...	

- Impact on inconsistent retrievals:
 - ▶ Fund transfer in V and branch total in W
 - ▶ If retrieval runs along update, then inconsistent
 - ▶ If retrieval is before/after update, then consistent

Transactions

Conflicting operations

Operations of different transactions		Conflict	Reason
<i>read</i>	<i>read</i>	No	Because the effect of a pair of <i>read</i> operations does not depend on the order in which they are executed
<i>read</i>	<i>write</i>	Yes	Because the effect of a <i>read</i> and a <i>write</i> operation depends on the order of their execution
<i>write</i>	<i>write</i>	Yes	Because the effect of a pair of <i>write</i> operations depends on the order of their execution

- Conflicting operations: the result depends on execution order
- Consider *read* and *write*
 - ▶ The effect includes value returned by *read* and value set by *write*.
 - ▶ *read* – *write* and *write* – *write* conflict
- Order of conflicting operations:
 - ▶ Serial equivalence of two transactions: conflicting operations are executed in the same order for both transactions.

Transactions

Conflicting operations

Transaction T :	Transaction U :
$x = \text{read}(i)$	
$\text{write}(i, 10)$	
	$y = \text{read}(j)$
	$\text{write}(j, 30)$
$\text{write}(j, 20)$	
	$z = \text{read}(i)$

- Transactions T and U :
 - ▶ $T : x = \text{read}(i); \text{write}(i, 10); \text{write}(j, 20);$
 - ▶ $U : y = \text{read}(j); \text{write}(j, 30); z = \text{read}(i);$
- Consider the interleaving in the figure.
- Not serially equivalent.
 - ▶ Swapped writes change value.
 - ▶ Fact remains unchanged.
- Serial Equivalence Conditions:
 - ▶ T accesses both i, j before U .
 - ▶ U accesses both i, j before T .

Transactions

Conflicting operations

- Concurrency Control Protocols: three approaches
 - ▶ Locking:
 - Locks are set on objects before access.
 - Only locked transactions can access the object.
 - May cause deadlocks.
 - ▶ Optimistic Concurrency Control
 - Free operation until commit.
 - Check conflicts before commit.
 - Abort conflicting transactions.
 - ▶ Timestamp Ordering
 - Transactions are timestamped.
 - Operations are compared with timestamps to decide execution order.
 - Delayed operations wait; rejected transactions are aborted.

Transactions

Recoverability from aborts

- Recoverability Requirement: log committed transactions.
 - ▶ Aborted transactions: no effect.
- Issues from aborts: dirty reads, premature writes.

Transactions

Recoverability from aborts - Dirty Reads

Transaction <i>T</i> :	Transaction <i>U</i> :
<i>a.getBalance()</i>	<i>a.getBalance()</i>
<i>a.setBalance(balance + 10)</i>	<i>a.setBalance(balance + 20)</i>
<i>balance = a.getBalance()</i> \$100	
<i>a.setBalance(balance + 10)</i> \$110	
	<i>balance = a.getBalance()</i> \$110
	<i>a.setBalance(balance + 20)</i> \$130
	<i>commit transaction</i>
<i>abort transaction</i>	

- Isolation property: uncommitted states should be invisible.
 - ▶ Dirty Reads: reading uncommitted changes.
- Example Scenario
 - ▶ *T* updates account A.
 - ▶ *U* reads account A.
 - ▶ *T* aborts after *U* commits (dirty read).
- *U*'s commit irreversible.

Transactions

Recoverability of transactions - Cascading aborts

- Recoverability
 - ▶ U delays commit until observed transactions (e.g. T) commit.
 - ▶ U aborts, if observed transaction aborts.
- Cascading Aborts:
 - ▶ If T aborts, U must abort too.
 - ▶ Other transactions seeing U must also abort.
 - ▶ To avoid cascading aborts, read from committed transactions only.
 - ▶ Read operations delayed until transactions commit/abort.

Transactions

Premature writes

Operations of the *Account* interface

deposit(amount)

deposit *amount* in the account

withdraw(amount)

withdraw *amount* from the account

getBalance() → *amount*

return the balance of the account

setBalance(amount)

set the balance of the account to *amount*

Operations of the *Branch* interface

create(name) → *account*

create a new account with a given name

lookUp(name) → *account*

return a reference to the account with the given name

branchTotal() → *amount*

return the total of all the balances at the branch

- Premature Writes: Aborts impact interaction of writes.
- T and U modify balance.

- If U aborts, T commits:
 - ▶ Balance should be \$105
 - ▶ Before image of aborting process (U) is restored
 - ▶ U 's before image \$105. OK
- If T aborts after U commits:
 - ▶ Balance should be \$110
 - ▶ The before image of aborting process (T) is restored (\$100)

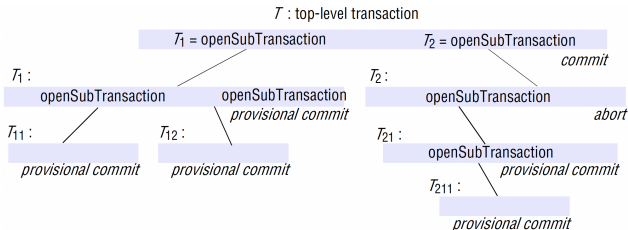
Transactions

Strict executions of transactions - Tentative versions

- Strict Executions:
 - ▶ Delay read/write operations
 - ▶ Prevents dirty reads and premature writes
 - ▶ Waits for commits/aborts
- Tentative Versions:
 - ▶ Removable updates
 - ▶ Volatile memory storage
 - ▶ Transfer upon commit

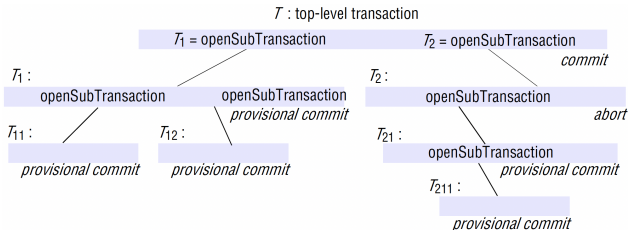
Transactions

Nested transactions until when



- Nested transactions modularity
- Top-level transaction
- Subtransactions:
 - ▶ T starts T_1 , T_2 ;
 - ▶ T_1 starts T_{11} , T_{12} ;
 - ▶ T_2 starts T_{21} , T_{211} .

Nested transactions



- Subtransactions are treated as atomic by parents
- Subtransactions can run concurrently.
- Subtransactions can fail independently.

Nested transactions

- Flat transactions
- Nested Transactions Advantages:
 - ▶ Concurrency benefits
 - ▶ Independent commit/abort

Nested transactions

- Committing Rules for Nested Transactions:
 - ▶ A transaction can only commit or abort after its child transactions are completed.
 - ▶ When a subtransaction completes, it makes an independent decision either to commit provisionally or to abort.
 - ▶ If a parent transaction aborts, all its subtransactions must abort.
 - ▶ A parent can choose to commit even if a subtransaction aborts.
 - ▶ If the top-level transaction commits, all provisionally committed subtransactions can also commit, unless their ancestors have aborted.
- Permanent Effects: Not permanent until top-level commit.

Nested transactions

- Top-level aborts on subtransaction abort.
- *Transfer* transaction requires both commits.

a.deposit(100)

b.withdraw(100)

- ▶ Both subtransactions needed.
- ▶ Abort of *withdraw* affects parent.
- ▶ Parent abort undoes all.

Locks

- Transactions require serial scheduling.
- Exclusive locks for serialization.
 - ▶ Server locks objects that a client is about to use.
 - ▶ Another requesting client waits for unlock.

Locks

Transaction <i>T</i> :		Transaction <i>U</i> :	
<i>balance</i> = <i>b.getBalance()</i> <i>b.setBalance</i> (<i>bal</i> *1.1) <i>a.withdraw</i> (<i>bal</i> /10)		<i>balance</i> = <i>b.getBalance()</i> <i>b.setBalance</i> (<i>bal</i> *1.1) <i>c.withdraw</i> (<i>bal</i> /10)	
Operations	Locks	Operations	Locks
<i>openTransaction</i>		<i>openTransaction</i>	
<i>bal</i> = <i>b.getBalance()</i>	lock <i>B</i>	<i>bal</i> = <i>b.getBalance()</i>	waits for <i>T</i> 's lock on <i>B</i>
<i>b.setBalance</i> (<i>bal</i> *1.1)		<i>openTransaction</i>	
<i>a.withdraw</i> (<i>bal</i> /10)	lock <i>A</i>	<i>bal</i> = <i>b.getBalance()</i>	waits for <i>T</i> 's lock on <i>B</i>
<i>closeTransaction</i>	unlock <i>A</i> , <i>B</i>	...	
			lock <i>B</i>
		<i>b.setBalance</i> (<i>bal</i> *1.1)	
		<i>c.withdraw</i> (<i>bal</i> /10)	lock <i>C</i>
		<i>closeTransaction</i>	unlock <i>B</i> , <i>C</i>

- Locking causes wait.
- Unlocking allows access.
- Serial equivalence required.

Locks

- No new locks after release of a lock.
- Growing phase for locks.
- Shrinking phase for release.
- Two-phase locking concept.

Locks

- Delay transactions until previous transactions abort or commit.
- Strict execution concept.
- Strict Two-Phase Locking: Locks until completion.
 - ▶ Locks prevent access.
- This way, locks persist until permanent storage...
 - ▶ Server holds many objects.
 - ▶ Transactions accessed a few, with rare conflicts.
 - ▶ Locks limit access.
 - ▶ Concurrency granularity matters.

Locks

- Efficient locking scheme needed.
- Many readers/single writer.
- Read lock before a read; write lock before a write.
- Concurrent read operations do not conflict.
- Multiple transactions can share a read lock.
- "Read lock" = "shared lock".

Locks

<i>For one object</i>		<i>Lock requested</i>	
		<i>read</i>	<i>write</i>
<i>Lock already set</i>	<i>none</i>	OK	OK
	<i>read</i>	OK	wait
	<i>write</i>	wait	wait

- Operation conflict rules:
 - ▶ Read prevents write.
 - ▶ Write prevents read/write.
- Write lock delayed by read lock.
- Read/write lock delayed by write lock.

Locks

- Inconsistent retrievals: Caused by conflicts between read and write operations without concurrency control.
- Prevention of inconsistent retrievals:
 - ▶ Retrieval before/after updates.
 - ▶ Read locks delay updates, if retrieval comes first.
 - ▶ Retrieval is delayed after update, if it comes second.
- Lost updates.
- Prevention of lost updates:
 - ▶ Set read locks and then promote them to write.
- Lock promotion:
 - ▶ Read lock cannot promote, if it is shared.
 - ▶ Request write lock; wait for other read locks to be released.

Locks

<i>For one object</i>		<i>Lock requested</i>	
		<i>read</i>	<i>write</i>
<i>Lock already set</i>	<i>none</i>	OK	OK
	<i>read</i>	OK	wait
	<i>write</i>	wait	wait

- Exclusivity of locks
 - ▶ Write locks exclusive
 - ▶ No other locks allowed
- Strict Two-Phase Locking Rules:
 - ▶ No client access to un/lock operations
 - ▶ Locking via read/write requests
 - ▶ Unlocking on commit or abort

Locks

Locking rules for nested transactions

- Sets of nested transactions are isolated.
- Transactions in a set are isolated.
 - ▶ No concurrent parent-child execution.
 - ▶ Parent retains lock transferring it to child temporarily while it executes.
 - ▶ Concurrent sibling executions are possible, if they serialize access.

Locks

Locking rules for nested transactions

- Lock rules for acquisition and release:
 - ▶ Acquire a read lock, if there is no active write lock.
 - ▶ Acquire a write lock, if there is no other locks.
 - ▶ Inherited locks from children.
 - ▶ Locks from aborted subtransactions are discarded unless retained by the parent.
- Lock acquisition between subtransactions at the same level is possible, if access is serialized.

Locks

Deadlocks


Deadlock with write locks

Transaction <i>T</i>		Transaction <i>U</i>	
Operations	Locks	Operations	Locks
<i>a.deposit(100);</i>	write lock <i>A</i>		
		<i>b.deposit(200)</i>	write lock <i>B</i>
<i>b.withdraw(100)</i>			
...	waits for <i>U</i> 's lock on <i>B</i>	<i>a.withdraw(200);</i>	waits for <i>T</i> 's lock on <i>A</i>
...		...	
...		...	

- Deadlock: mutual waiting of two transactions.
- E.g. Deposit and withdraw both acquire write locks.
- *T* and *U* are both blocked on locked accounts.

Locks


Definition of Deadlock



- Deadlocks can be identified with wait-for graphs
- Nodes are transactions, edges are wait-for relationships.
- Cycles in wait-for graphs indicate deadlocks.
- Breaking a deadlock requires aborting a transaction in the cycle.

Locks

Definition of Deadlock



- Multiple cycles possible:
 - ▶ A transactions can be in multiple cycles.
 - ▶ Aborting such transactions breaks all involved cycles.

Locks

Deadline prevention - Upgrade locks

- To prevent deadlocks:
 - ▶ One can lock all objects used by a transaction initially.
 - ▶ This restricts resource access.
 - ▶ Predicting necessary locks is difficult.

Locks

Deadlock detection

- Deadlocks: cycles in wait-for graph.
- Abort a transaction to resolve the deadlock.
- Lock Manager's Role in Wait-for Graph Maintenance: adding removing edges based on setLock and unLock operations
- Cycle Checking Frequency:
 - ▶ Check cycles every time an edge is added or less frequently.
- Transaction abortion and wait-for graph maintenance:
 - ▶ Abort a transaction on deadlock.
 - ▶ Remove aborted transaction's node and edges.
 - ▶ Choosing the transaction to abort is difficult.

Locks

Timeouts

- Timeouts is one way to handle deadlocks.
- A Lock has a limited time during which it is unbreakable.
 - ▶ When timeout expires, lock becomes vulnerable.
 - ▶ If there is no competition, the lock remains.
 - ▶ With competition, a transaction is aborted.
- Problems with Timeouts:
 - ▶ Abortions are possible even without deadlock.
 - ▶ Difficulty in determining timeout periods.
- Distributed Deadlocks:
 - ▶ Objects across multiple servers complicate deadlock detection.

Locks

Increasing concurrency in locking schemes

- Concurrency in Locking:
 - ▶ Increase concurrency despite conflicts.
 - ▶ Two approaches:
 - Two-Version Locking: Delay exclusive locks.
 - Hierarchic Locks: Mixed-granularity locks.

Locks

Two version locking

- Optimistic transaction management.
- Higher concurrency potential.
- Types of locks
 - ▶ Read Lock: Multiple read locks possible.
 - ▶ Write Lock: Single write lock at a time.
 - ▶ Commit Lock: Final approval lock.
- Read delayed by commit.

Locks

Two version locking

<i>For one object</i>		<i>Lock to be set</i>		
		<i>read</i>	<i>write</i>	<i>commit</i>
<i>Lock already set</i>	<i>none</i>	OK	OK	OK
	<i>read</i>	OK	OK	wait
	<i>write</i>	OK	wait	–
	<i>commit</i>	wait	wait	–

- Rules:
 - ▶ Read lock granted unless commit lock exists.
 - ▶ Write lock granted unless write/commit lock exists.
 - ▶ Commit attempts to convert locks.

Locks

Hierarchic locks

- Suitable for application which require different locking granularities.
 - ▶ E.g. some operations require lock at account level or others at branch level.
- Advantages: Mixed granularity reduces overhead and fewer locks needed.
- Disadvantages: The compatibility tables and the rules for promoting locks are more complex.

Optimistic concurrency control

- Disadvantages of Locking:
 - ▶ Lock maintenance adds overhead.
 - ▶ Locks can lead to deadlocks.
Cascading aborts.
- Alternative optimistic approach: Low conflict assumption.
- Transactions proceed until a closeTransaction request.
- If a conflict arises at the time of a closeTransaction request, a transaction is aborted and will need to be restarted.

Optimistic concurrency control

- Transaction Phases:
 - ▶ Working Phase:
 - Each transaction uses a tentative version of the object.
 - Write operation creates tentative values.
 - Read operation is performed immediately.
 - ▶ Validation Phase:
 - Upon a closeTransaction request, the transaction is validated for conflicts.
 - Successful validation allows commitment.
 - ▶ Update Phase:
 - Tentative changes become permanent.
 - Read-only transactions commit immediately.
 - Write transactions commit after recording tentative versions in permanent storage.

Optimistic concurrency control

Validation of transactions

- Validation is based on read-write conflict rules
 - ▶ Conflicts can occur between *overlapping transactions*.
 - ▶ A transaction's overlapping transactions are those that did not commit at the time that transaction started.
- Unique transaction numbering upon closeTransaction request
 - ▶ T_i precedes T_j , if $i < j$.

Optimistic concurrency control

Validation of transactions

- Rules for validating a transaction T_v

T_v	T_i	<i>Rule</i>	
<i>write</i>	<i>read</i>	1.	T_i must not read objects written by T_v .
<i>read</i>	<i>write</i>	2.	T_v must not read objects written by T_i .
<i>write</i>	<i>write</i>	3.	T_i must not write objects written by T_v and T_v must not write objects written by T_i .

Optimistic concurrency control

Validation of transactions

- Single transaction validation at a time
 - ▶ No overlap in update phase
- Validation Types:
 - ▶ Backward Validation against preceding overlapping transactions.
 - ▶ Forward Validation against succeeding overlapping transactions.

Optimistic concurrency control

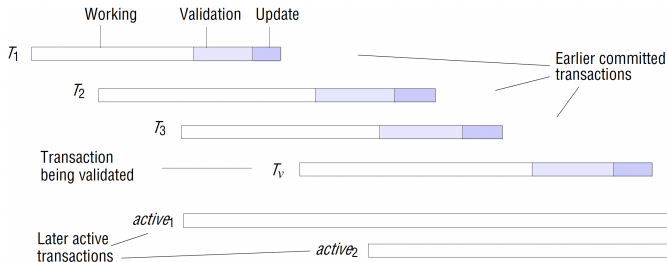
Backward validation

- Backward validation validates transactions after they read data.
- If "startTn" and "finishTn" are largest transaction numbers at the start of working and validation phases, then

```
boolean valid = true;  
for (int  $T_i = startTn + 1$ ;  $T_i \leq finishTn$ ;  $T_i++$ ) {  
    if (read set of  $T_v$  intersects write set of  $T_i$ ) valid = false;  
}
```

Optimistic concurrency control

Backward validation



- T_v to be validated.
- Earlier committed transactions: T_1 , T_2 , T_3
- T_1 committed before T_v started its working phase.
- Read set of T_v must be compared with the write sets of T_2 and T_3 .

Optimistic concurrency control

Backward validation

- Failed validation aborts transaction
- Write sets of old transactions need to be retained until all overlapping transactions validate.

Optimistic concurrency control

Forward validation

- Forward validation compares the write set of T_v with the read sets of overlapping active transactions.

```
boolean valid = true;  
for (int  $T_{id} = active_1$ ;  $T_{id} \leq active_N$ ;  $T_{id}++$ ) {  
    if (write set of  $T_v$  intersects read set of  $T_{id}$ ) valid = false;  
}
```

- Compare T_v with active transactions $active_1$ to $active_N$.

Optimistic concurrency control

Forward validation

- Conflict Resolution Options:
 - ▶ Abort Validating Transaction: Simple, possible unnecessary aborts.
 - ▶ Abort Conflicting Transactions: Commit T_v , abort conflicting active transactions.
 - ▶ Defer Validation: Wait for conflicts to be resolved.

Optimistic concurrency control

Comparison of forward and backward validation

- Forward vs. backward validation:
 - ▶ Forward: flexible conflict resolution.
 - ▶ Read sets larger than write sets.
 - Backward: large read sets vs. old write sets.
 - Forward: small write sets vs. active read sets.

Discussion topic

A server manages the objects a_1, a_2, \dots, a_n . The server provides two operations for its clients:

read(i) returns the value of a_i ;

write($i, Value$) assigns $Value$ to a_i .

The transactions T and U are defined as follows:

T : $x = \text{read}(j)$; $y = \text{read}(i)$; $\text{write}(j, 44)$; $\text{write}(i, 33)$;

U : $x = \text{read}(k)$; $\text{write}(i, 55)$; $y = \text{read}(j)$; $\text{write}(k, 66)$.

Give two serially equivalent interleavings of the transactions T and U .