



# Project and written exam (1/2)

- 3 group tasks during the course (two thirds of the grade)
  - 2 small ones at the beginning: DONE!
  - 1 big task at the end implementing the architecture
    - Using **any technology**, not only the ones I'll introduce you
    - Containing **at least one monolithic and one distributed architecture**
    - Deadline: before the written exam (**January 13th**)!
- Written exam at the end of the course (one third of the grade)
  - **Only after submitting and passing** the tasks/project
  - Example from last year in Moodle
  - About **only all** the in presence lectures



# Project and written exam (2/2)

- You will have to deliver at the end of the project
  - A GitHub repository containing the **code**
  - The **history of the repo** should make clear the contribution of the different students in the group
  - It must contain a **description** (pdf, markdown, ...) of the structure of the architecture with references to the code and **how to run it**
    - Docker (file or compose) or Kubernetes (minikube)



- Schedule for the next lectures:
  - Nov. 14th in Lab. 3
  - (Nov. 16th Google DevFest 24)
  - Tuesday Nov. 19th at 8:45 in Aula C
  - Nov. 28th invited lecture by Unox
    - [https://www.unox.com/us\\_us/](https://www.unox.com/us_us/)
  - Dec. 5th invited lecture by Gianluca Caiazza
    - Zamperla project
    - <https://www.zamperla.com/>



## Task 2

Please submit a pdf file where you define **what architecture characteristics** among the ones we identified are **of interest** to your IT system, **relating** each architecture characteristics to the points of the kata that led you to identify that characteristic. For each architecture characteristic you picked up, please **"quantify"** exactly at what level you need it for your IT system.

PS: the pdf can contain also an **updated version of the kata** provided in the previous tasks (e.g., to number the various requirements and refer to them in the new task)



Ca' Foscari  
University  
of Venice

# Distributed architectures [Concepts]

Software architectures

Pietro Ferrara

[pietro.ferrara@unive.it](mailto:pietro.ferrara@unive.it)



# From monolithic to distributed

## Monolithic

- Initially, a unique ball of mud
- Then distinct layers or functionalities
  - n-tier
  - pipeline
  - microkernel
- Conceptually, in the same machine
- Practically, might be more than one
  - But the system is seen as a unique block of stone
- Monolithic architectures will be also replicated in distributed systems

## Distributed

- Initially, distinct networked computers
  - Communicating through messages
  - E.g., client-server
- Then, several machines communicating to provide a service
  - Service-based/oriented architecture
  - Event-based architecture
  - Microservices
- More complex architectures
- Network communications take time
- Physical distinction of services
- Potentially, many machines for one task



# API

- Application Programming Interface
  - How sw components communicate
- E.g., interface of Java classes
  - Public methods and fields
- Distributed systems more complex
  - I ping a machine and what I get back?
    - “Hello world!”
    - `<xml> ... </xml>`
    - `<html> ... </html>`
    - `{ "professor": { "name": "Pietro" } }`
  - Many possible choices
- From an architectural perspective...
  - We don't care!
  - But we need to exchange information!

## Java® Platform, Standard Edition & Java Development Kit Version 17 API Specification

This document is divided into two sections:

### Java SE

The Java Platform, Standard Edition (Java SE) APIs define the core Java platform for general-purpose computing. These APIs are in modules whose names start with `java`.

<https://docs.oracle.com/en/java/javase/17/docs/api/>

Some (very) good material on those topics  
(from a course of our bachelor in CS)





# REST communication

- A style, not a standard!
- Already found: RestController
  - A controller with a body
  - Plain text, XML, JSON, ...
- REST == Representational state transfer
- Stateless protocol
  - No session state (i.e., identifier)
  - All relevant information transferred through the messages
- Advantages: services are independent
  - Do not need to access session information
- Drawbacks: network bandwidth
  - Need to transfer a lot of infos

**Representational state transfer (REST)** is a [software architectural style](#) that describes a uniform [interface](#) between physically separate components, often across the [Internet](#) in a [client-server](#) architecture. REST defines four interface constraints:

- Identification of resources
- Manipulation of resources
- Self-descriptive messages and
- [Hypermedia as the engine of application state](#)<sup>[1]</sup>

[https://en.wikipedia.org/wiki/Representational\\_state\\_transfer](https://en.wikipedia.org/wiki/Representational_state_transfer)


A **stateless protocol** is a [communication protocol](#) in which the receiver must not retain [session](#) state from previous requests. The sender transfers relevant session state to the receiver in such a way that every request can be understood in isolation, that is without [reference](#) to session state from previous requests retained by the receiver.<sup>[1]</sup>

[https://en.wikipedia.org/wiki/Stateless\\_protocol](https://en.wikipedia.org/wiki/Stateless_protocol)



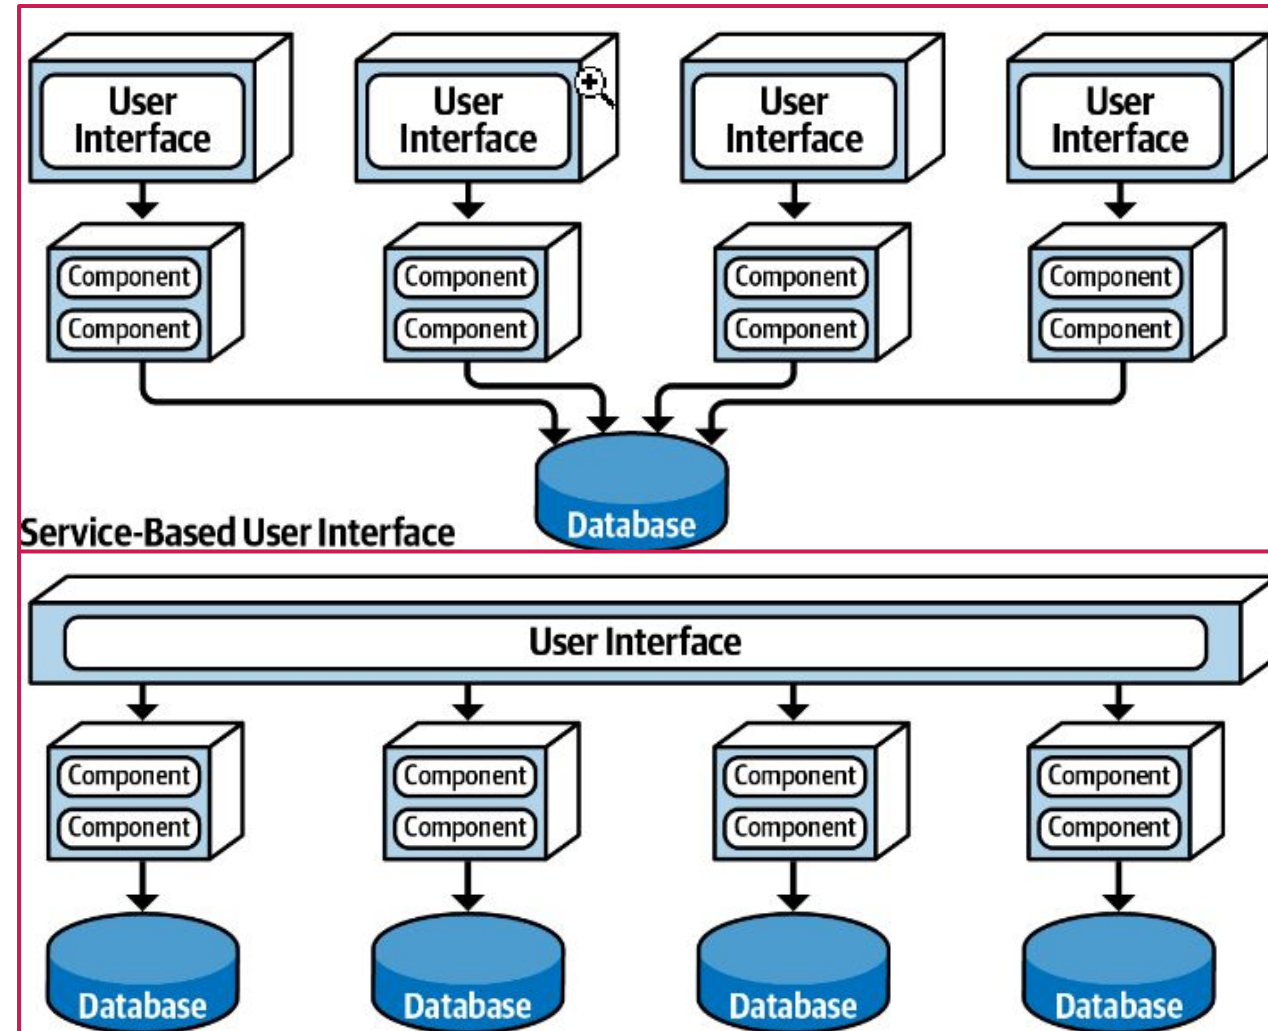
# Service-based architecture style

- The back-end comprises all the complex logic of the application
- It's usually the bottleneck
- Simple idea: split it into many services
  - Domain-based splitting
  - Still relying on a n-tier architecture
- A service is a portion of the application
  - Independent deployable unit
- The user interface collects data from some/all the different services
- Data is stored in a unique database



- Service-based user-interface
  - Each UI interacts with only one service
  - Require to access each service in isolation from the others
- Service-based database
  - Tables belong to one service
  - No relationship between different DBs
- An API layer can be added between user interface and services
  - Proxy knowing who provides what
- More components, more complexity
  - More scalability!

## Variations



# Your wooclap poll will be displayed here



Install the **Chrome** or  
**Firefox extension**



Make sure you are in  
**presentation mode**

**wooclap**



Ca' Foscari  
University  
of Venice

---



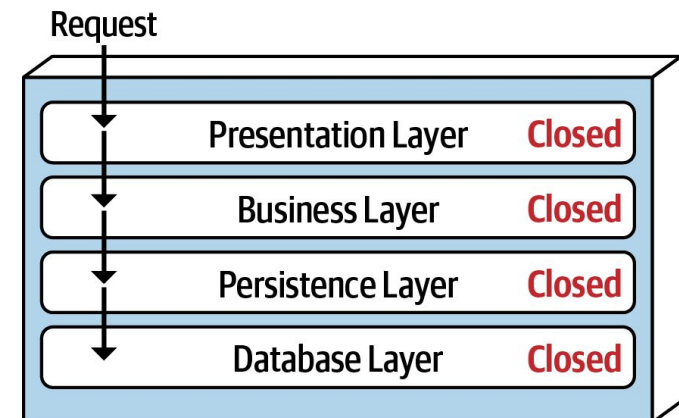

# Service-based architecture style

## Textbook solution

<b>Deployability</b>	★★★★	<b>Performance</b>	★★★
<b>Elasticity</b>	★★	<b>Reliability</b>	★★★★★
<b>Evolutionary</b>	★★★	<b>Scalability</b>	★★★
<b>Fault tolerance</b>	★★★★★	<b>Simplicity</b>	★★★
<b>Modularity</b>	★★★★★	<b>Testability</b>	★★★★★
<b>Overall cost</b>	★★★★★		

# Synchronous vs asynchronous


- So far, only synchronous components
  - I call a method: I wait for the returned value
  - I run a pipeline: each step waits for the results of the previous step
  - I rely on a distributed service: I make a request, I wait until I get a reply
- Obviously, this limits our architecture
  - If a component fails, everything fails
  - If a component is slow, everything is slow
  - If a component does not scale up, the system is not scalable
- We need to move to asynchronous
  - More complexity coming...




# Asynchronous capabilities

If we are posting  
a comment, no  
need to wait  
everything is  
finalized!

## Synchronous



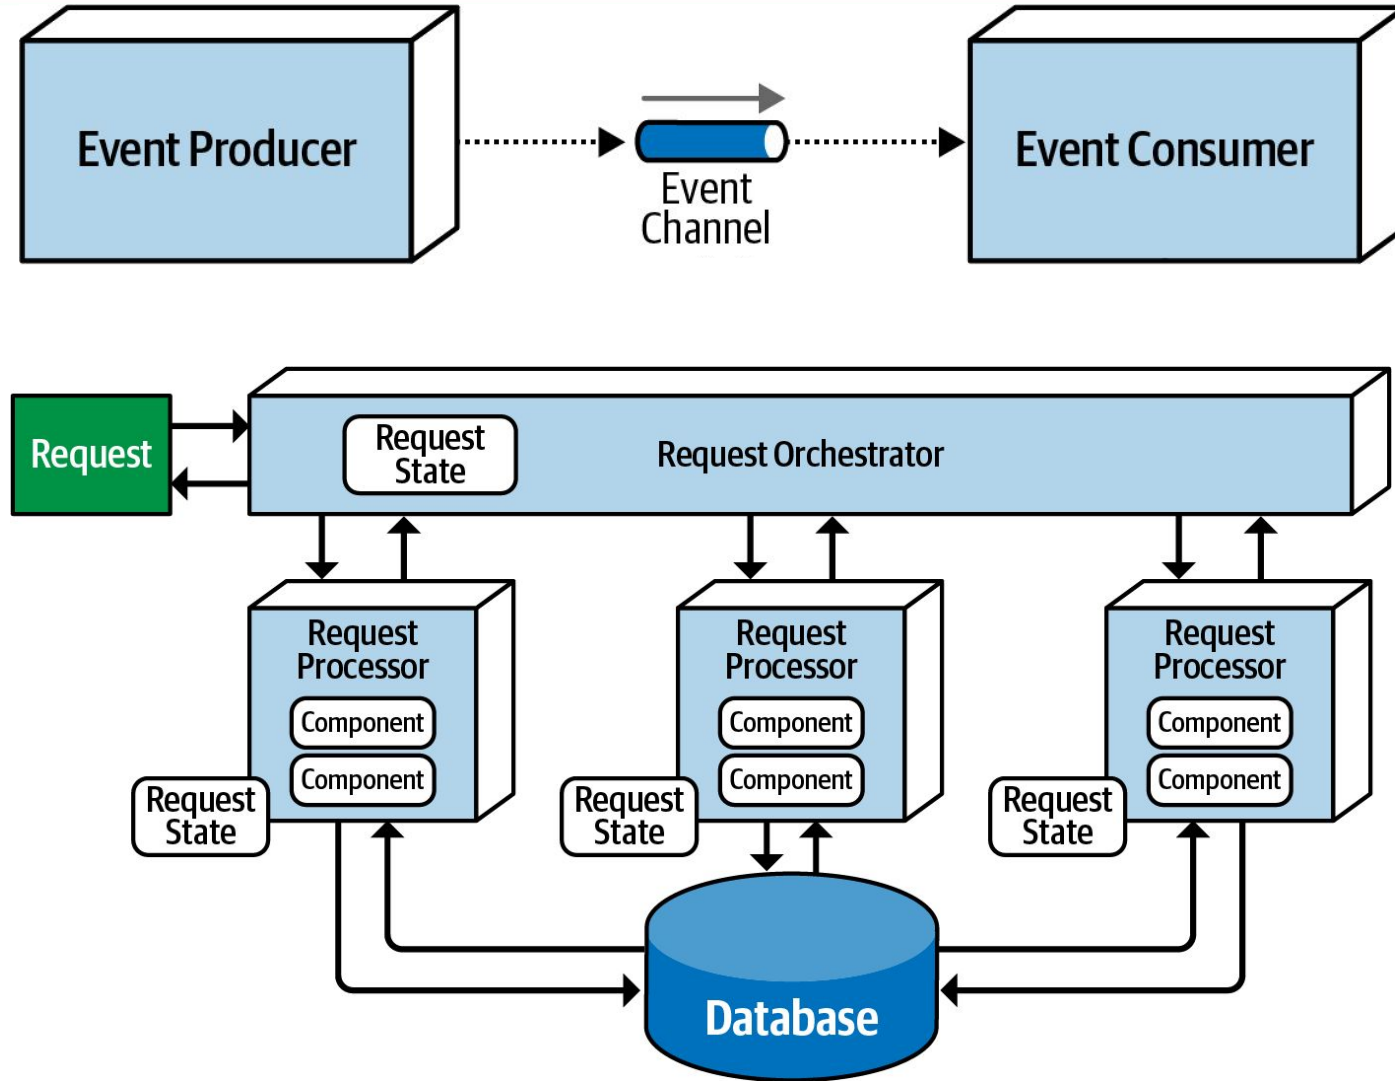
## Asynchronous





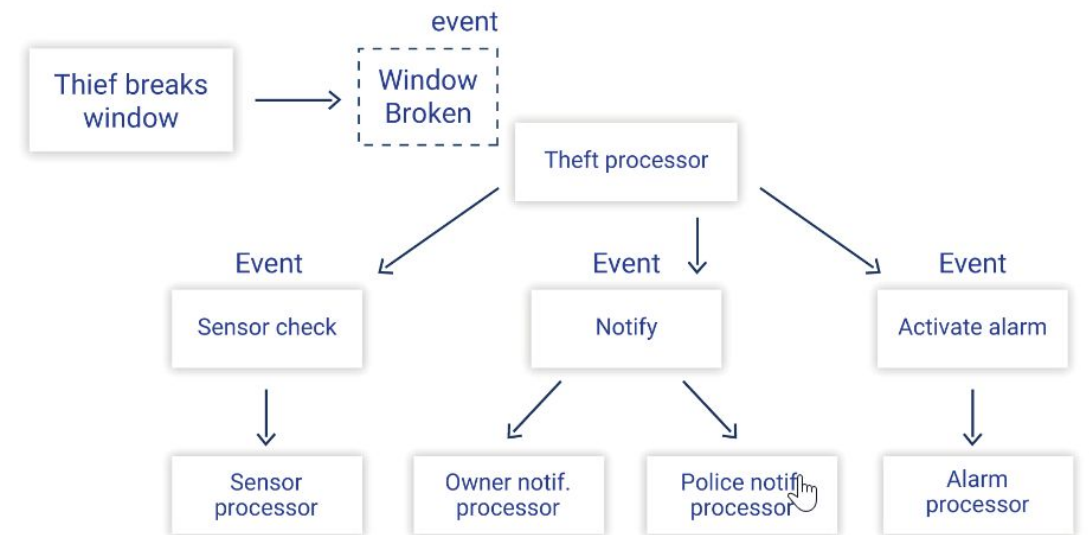
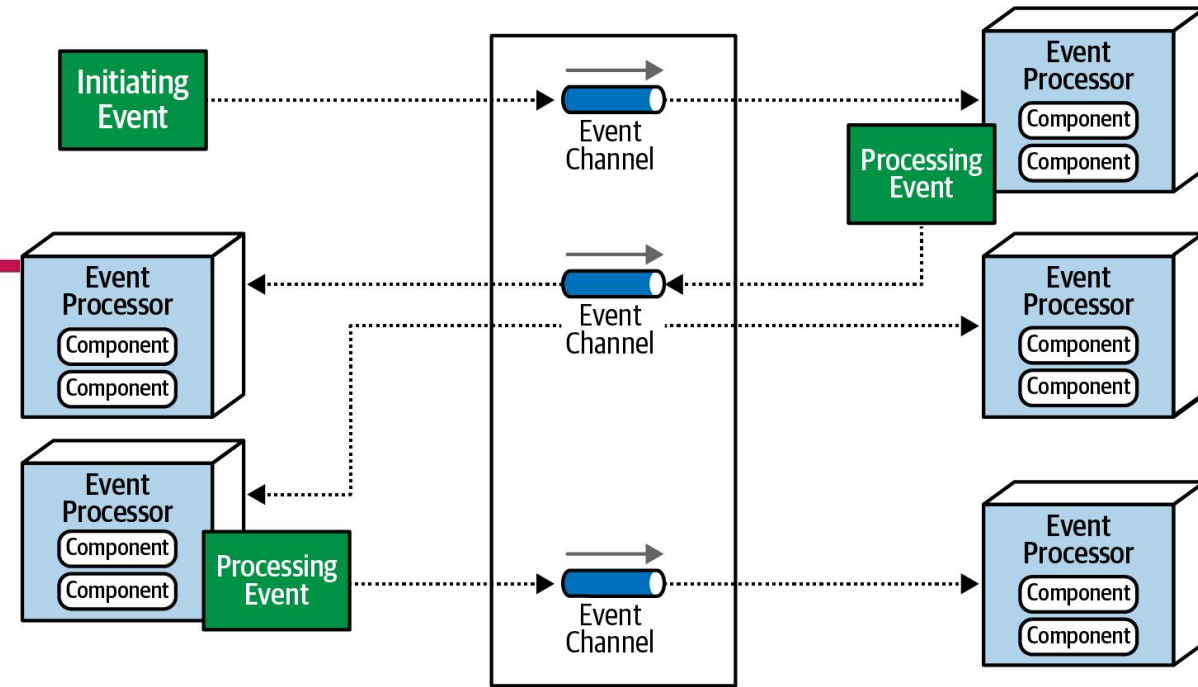
# Event-Driven Architecture style

- Request-response (APIs):
  - I ask for something (request)
  - I idle until I get a reply (response)
  - Method calls, RPCs, RESTs, ...
- Event-driven:
  - I ask for something
  - I do something else until
  - I get pinged: I have a reply!
- Three major components:
  - A producer of events
  - A communication (event) channel
  - A consumer of events
- Just the basic block...

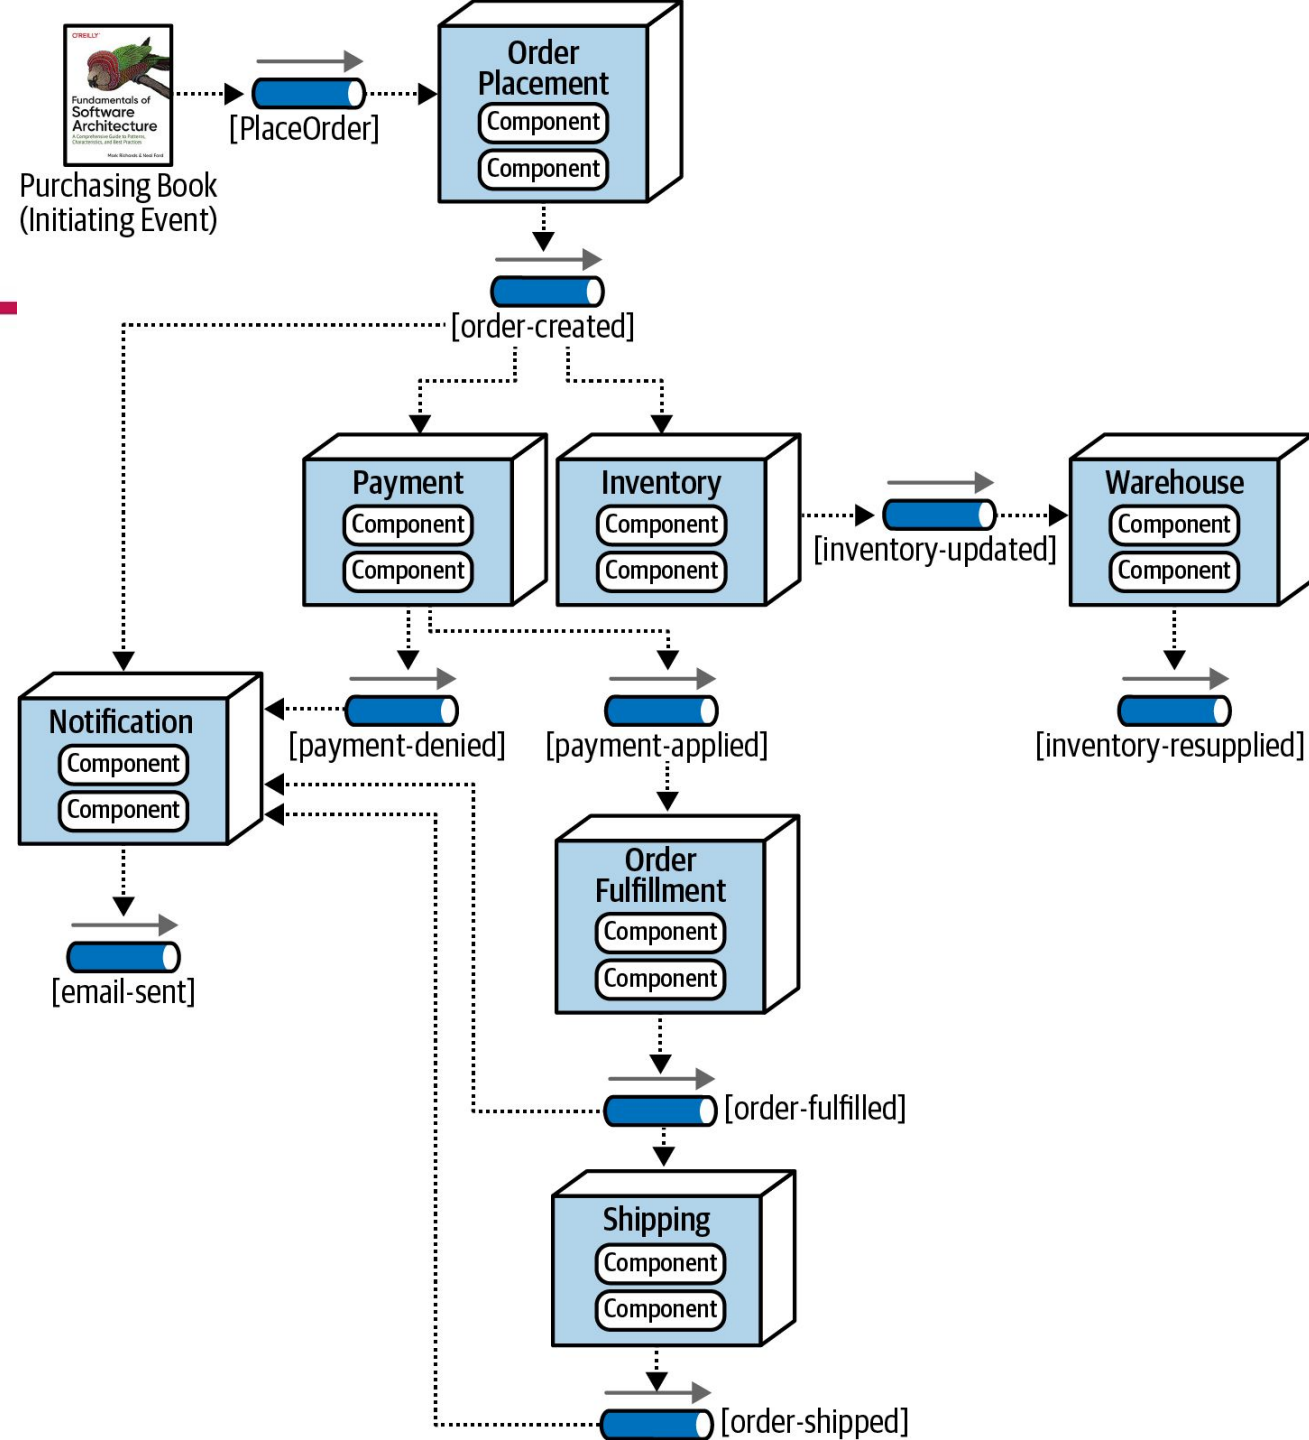


# Broker topology

- Messages flow among several processors
  - ... pipeline architecture?
- Initiating event: the begin of everything
  - Entrypoint!
- Event broker: communication channel
- Event processor: receives the event
- Processing event: the output of the event processor at the end of the processing
  - Another processor will take care of it
- Usually, publish-and-subscribe communication model through topics
  - RabbitMQ, Kafka, etc...

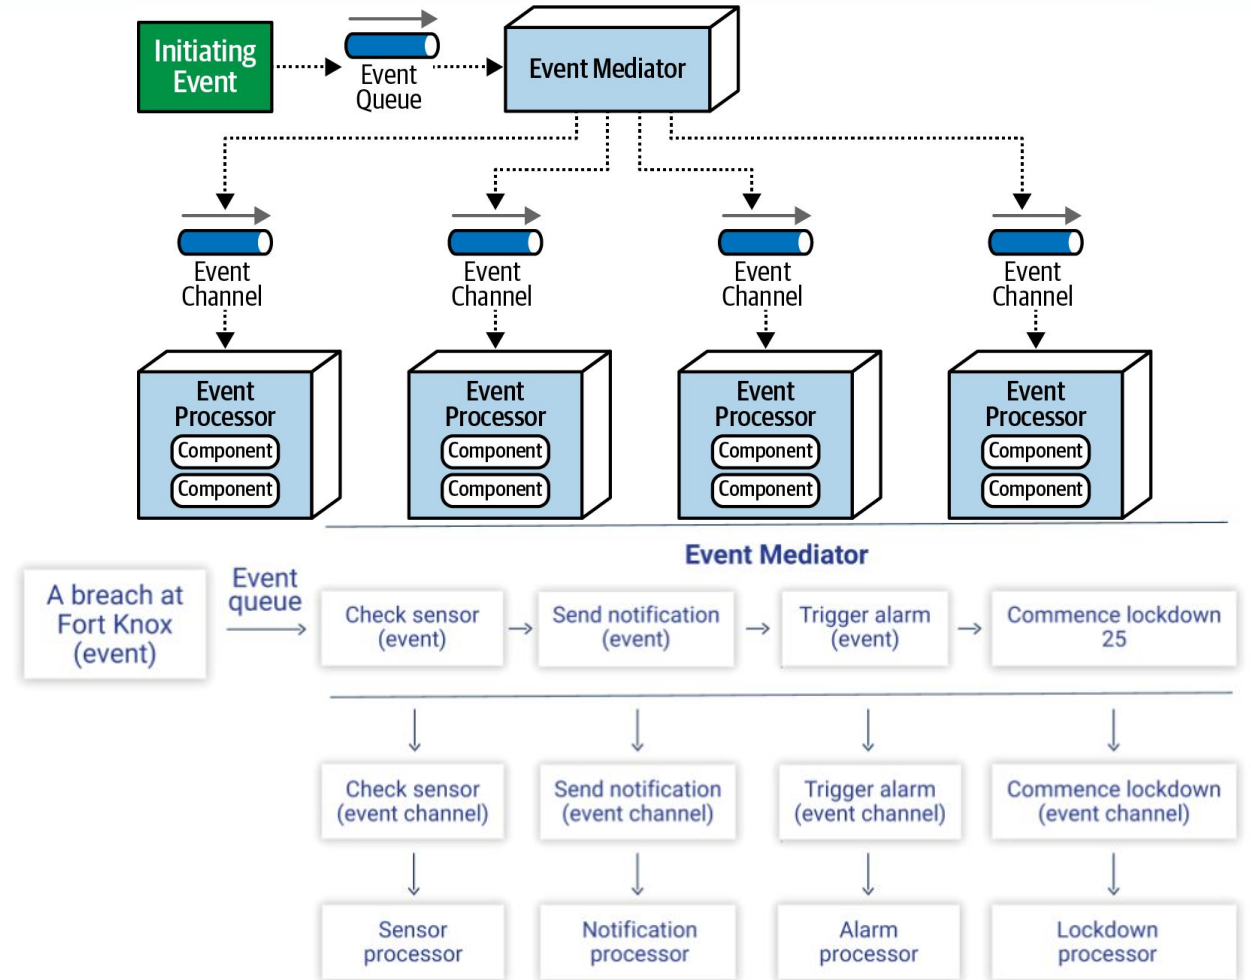


# An example of the broker topology



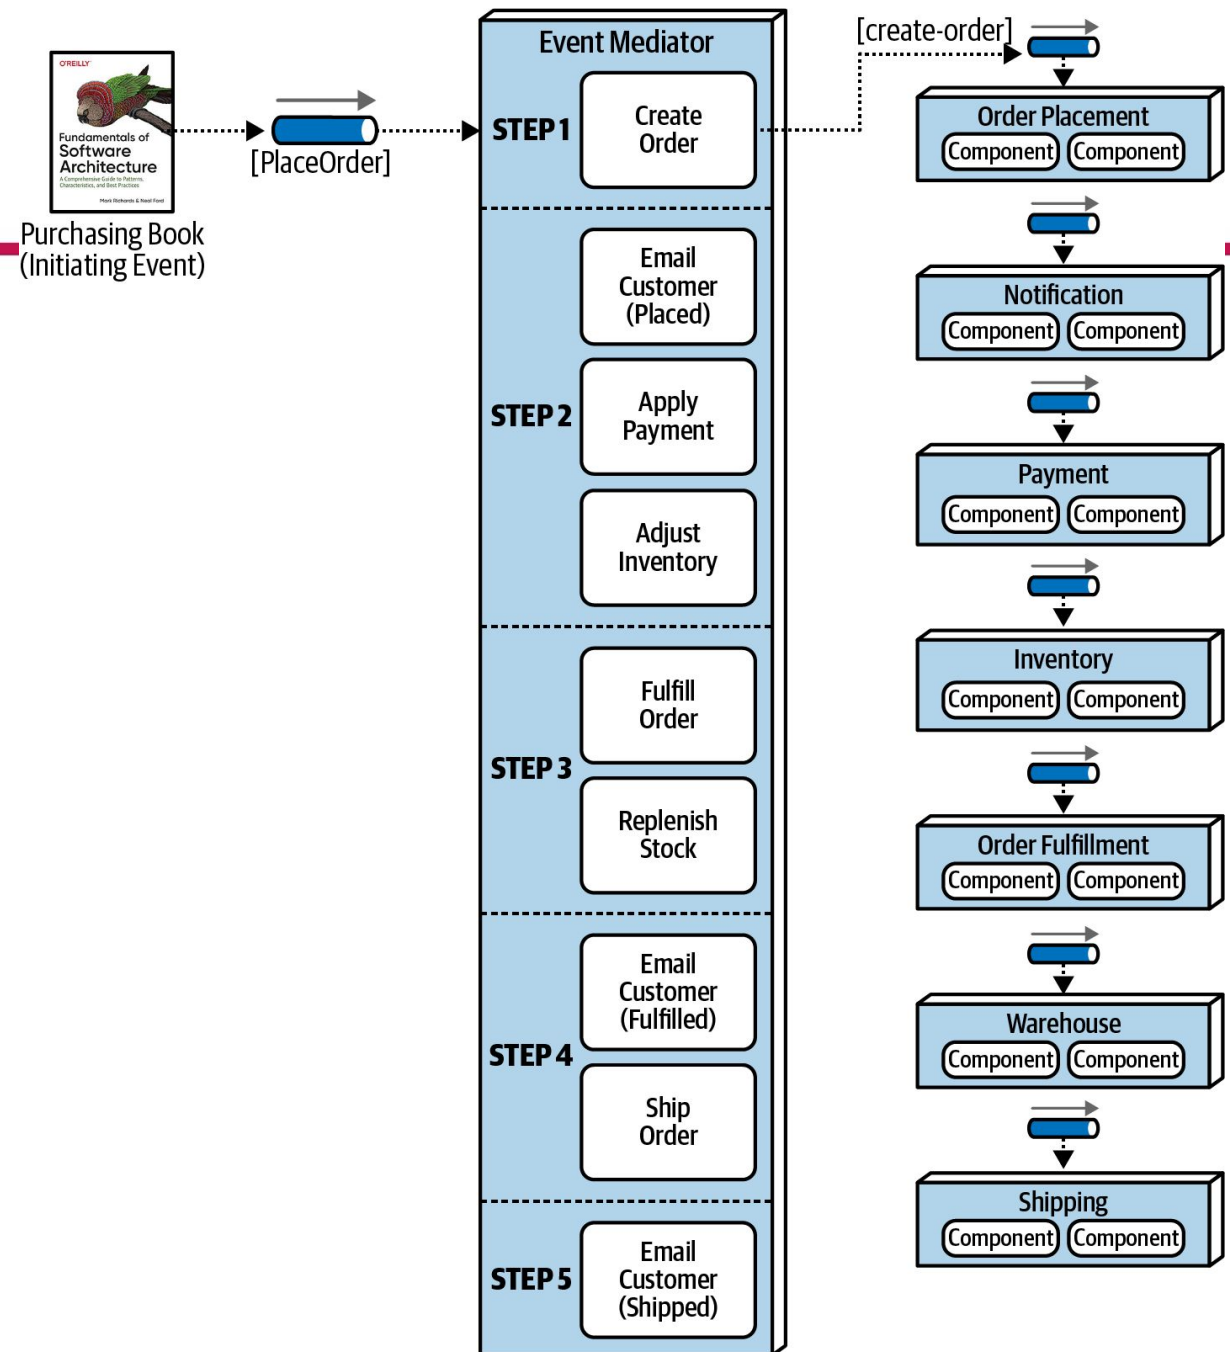
# Mediator topology

- Central mediator
  - Receive the original event
    - Aka request
  - Send events to all the other services involved
  - Collect the results of the process
  - Eventually gives back some info
- Multiple mediators
  - Communications by a queue with many mediators listening
- In some ways, it recalls request-response pattern
  - But in an asynchronous way!





# An example of the mediator topology







# Broker vs mediator

## Broker

### Advantages

- Highly decoupled event processors
- High scalability
- High responsiveness
- High performance
- High fault tolerance

### Disadvantages

- Workflow control
- Error handling
- Recoverability
- Restart capabilities
- Data inconsistency

## Mediator

### Advantages

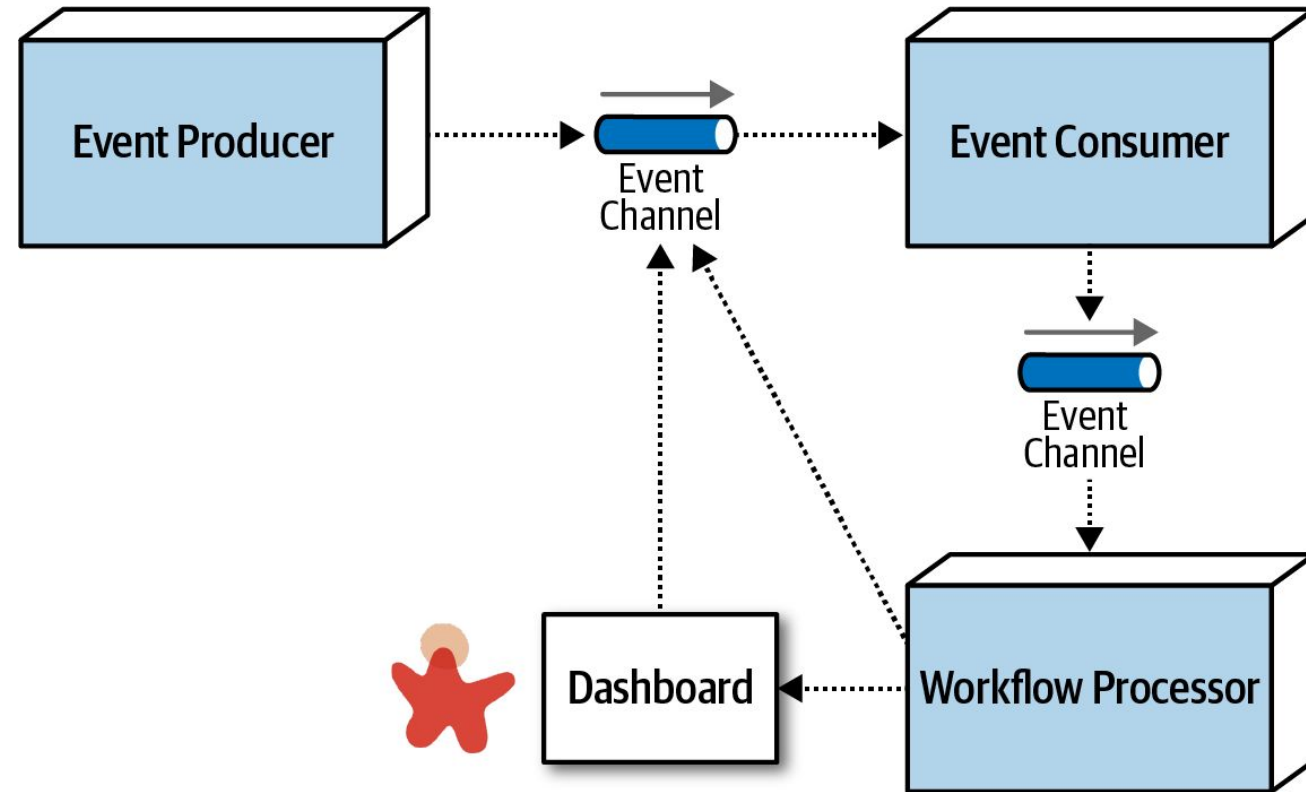
- Workflow control
- Error handling
- Recoverability
- Restart capabilities
- Better data consistency

### Disadvantages

- More coupling of event processors
- Lower scalability
- Lower performance
- Lower fault tolerance
- Modeling complex workflows

# Error handling (aka fault tolerance)

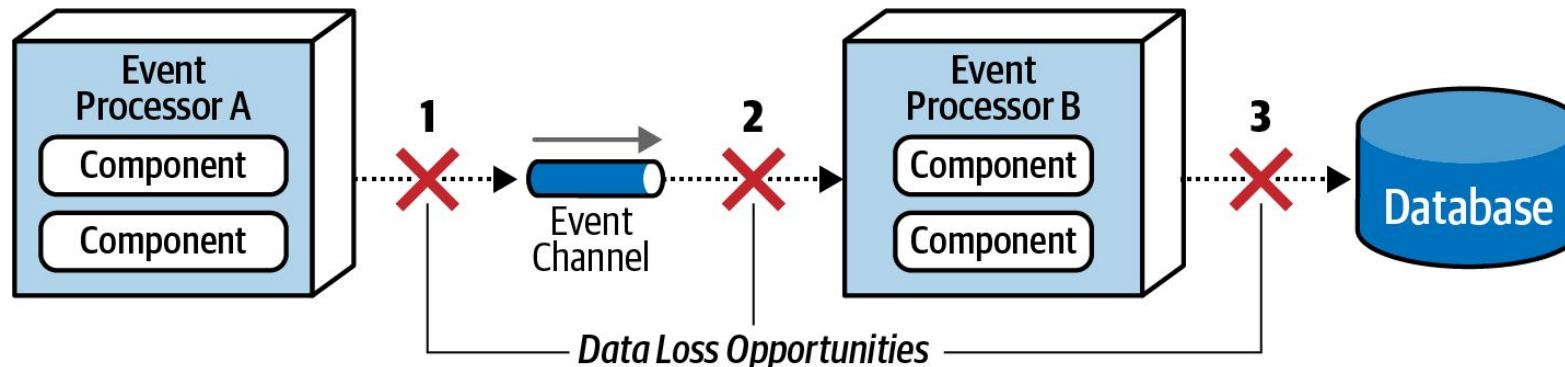
- Workflow processor
  - Check if a processor produced an error
  - If so, it potentially fixes the message and puts it back to the event channel
- The workflow processor automatically changes the data
  - Trying to repair it in some ways
  - Without any human intervention
- If it is not able to fix the message
  - Ask a user to give a look to it
  - Through some dashboard
    - Kind of email client





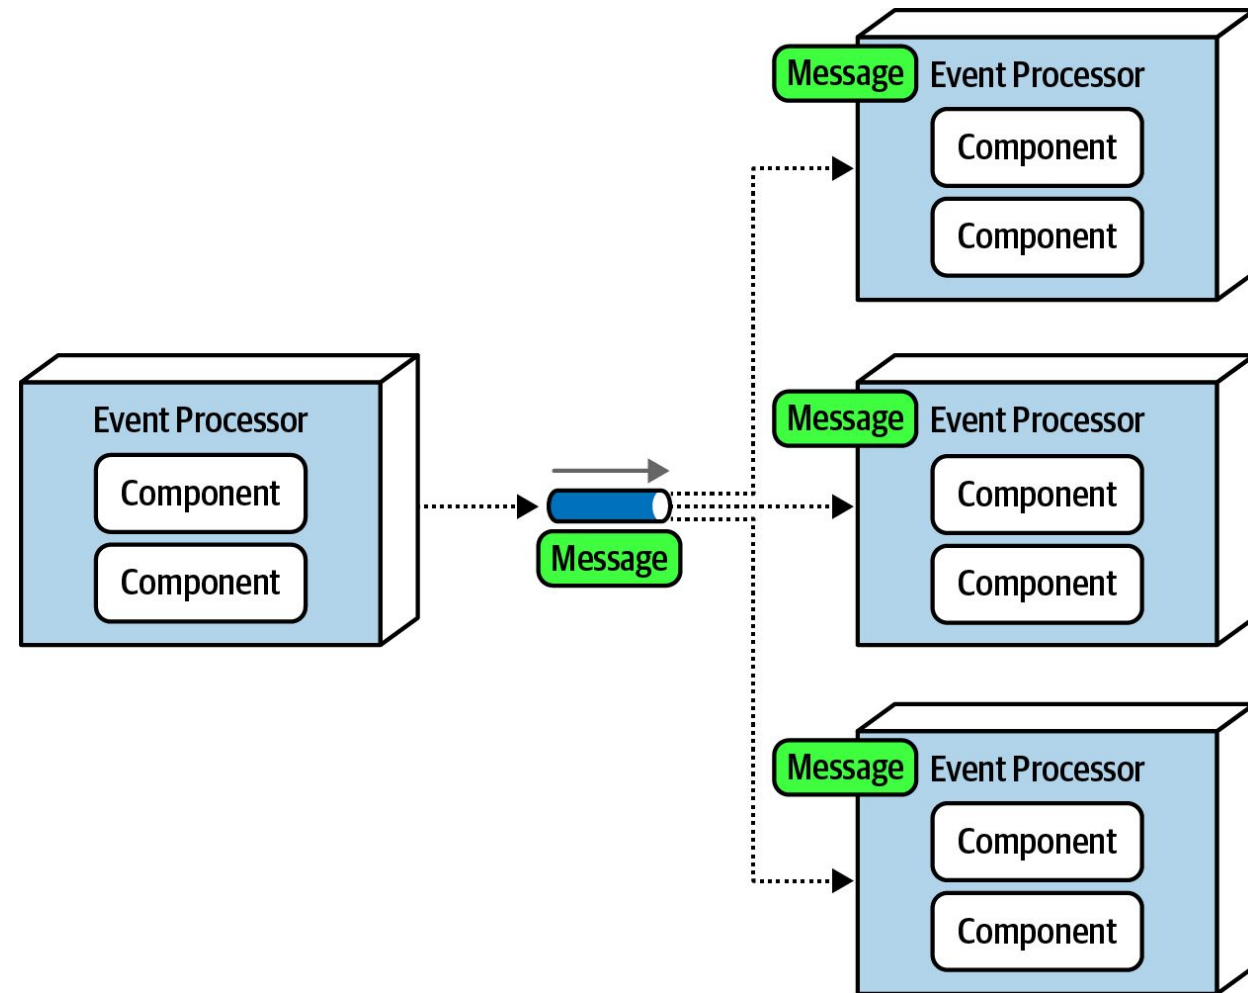
# Losing data

- Messages might be lost because
  - communication fails
  - an error of the event channel
- Rely on standard messaging techniques
  - Case 1 (a message never arrives to the queue): persistent message queues + synchronous send (guaranteed delivery)
  - Case 2 (a message is dequeued but the processor fails before processing it): client acknowledge mode (keep the message with the client ID, if fail then re-queue it)
  - Case 3 (unable to persist the message): ACID transactions via a database commit



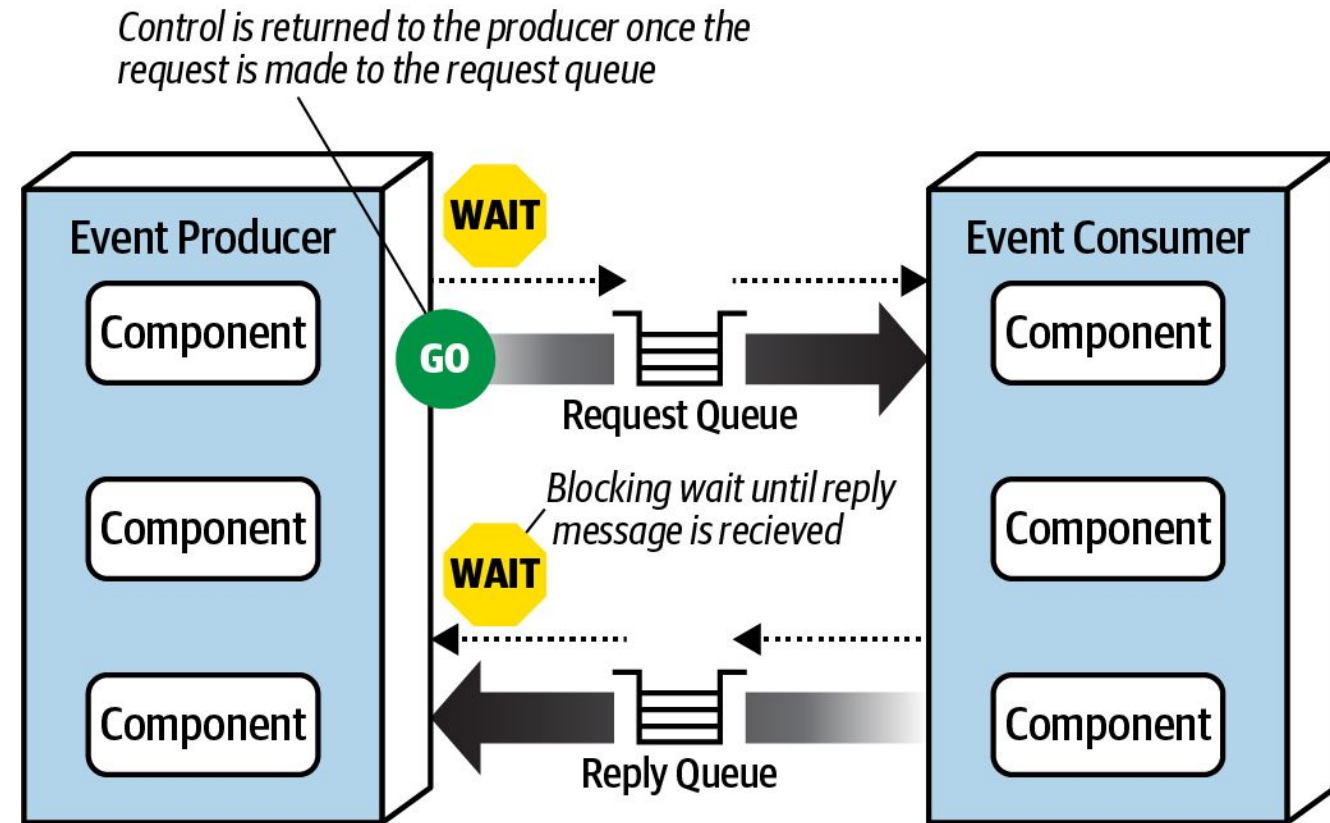
# Broadcast capabilities

- Event-based architecture does not know who process what!
  - Many subscribers might process the same message
- Highest level of decoupling between producer and consumer
- This enhances a lot
  - Scalability
  - Elasticity



# Request-Reply

- Everything asynchronous so far
- But sometimes we need to end up an action before moving on
- Adopt request-reply messaging
  - A queue receives the requests
  - After their processing, a message is sent to the reply queue
  - Then the producers will know that the processor finalized the task
- Quite more complex than standard REST request-response
- But we can do something else while waiting for the reply...



# Your wooclap poll will be displayed here



Install the **Chrome** or  
**Firefox extension**



Make sure you are in  
**presentation mode**

**wooclap**



Ca' Foscari  
University  
of Venice

---





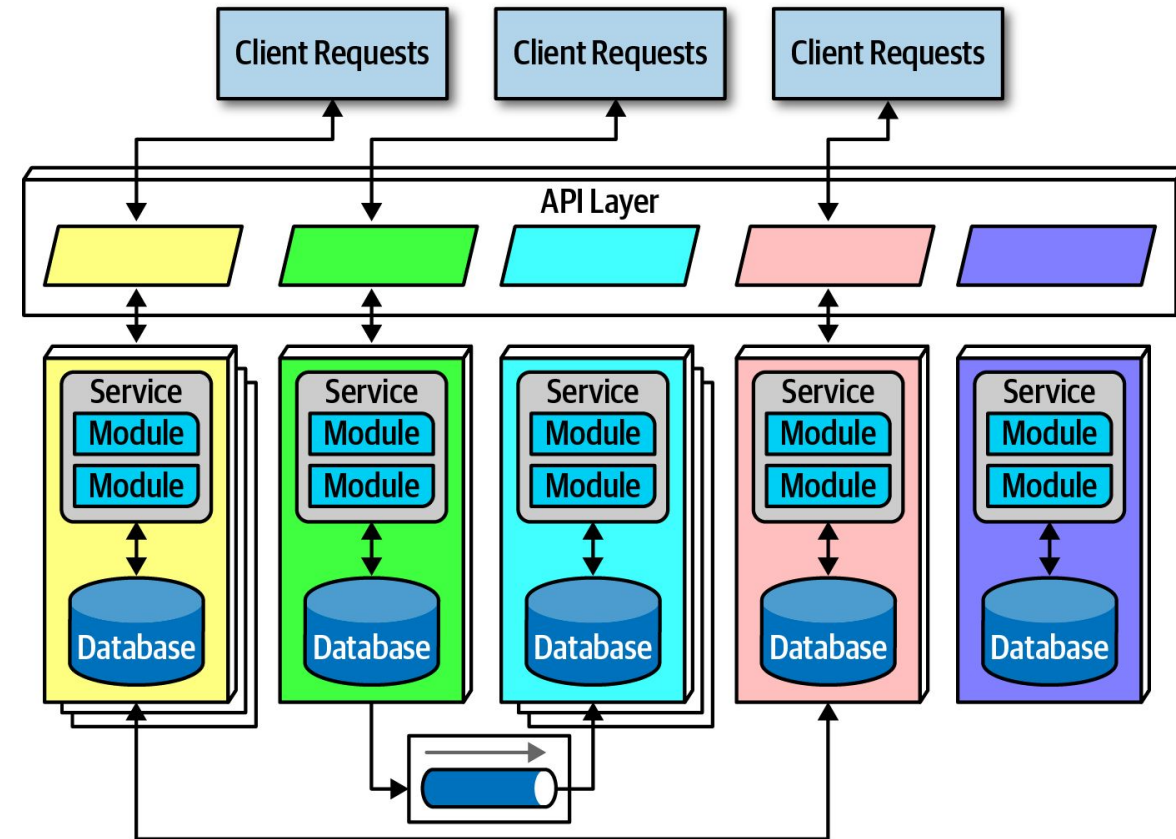
# Event-based architecture style

## Textbook solution

<b>Deployability</b>	★★★	<b>Performance</b>	★★★★★
<b>Elasticity</b>	★★★	<b>Reliability</b>	★★★
<b>Evolutionary</b>	★★★★★	<b>Scalability</b>	★★★★★
<b>Fault tolerance</b>	★★★★★	<b>Simplicity</b>	★
<b>Modularity</b>	★★★★★	<b>Testability</b>	★★
<b>Overall cost</b>	★★★		

# Microservices

- “Recent” big hype (for how long?)
  - Name appeared in 2014
- Derived from a common pattern
- Domain-driven design, bounded context
  - Each service has knowledge only about the part of its interest
    - No reusable classes/linked DBs in different services
- High level of decoupling, less code reuse
- Services are single-purpose (aka, micro)
- Each service includes all the necessary parts to operate independently





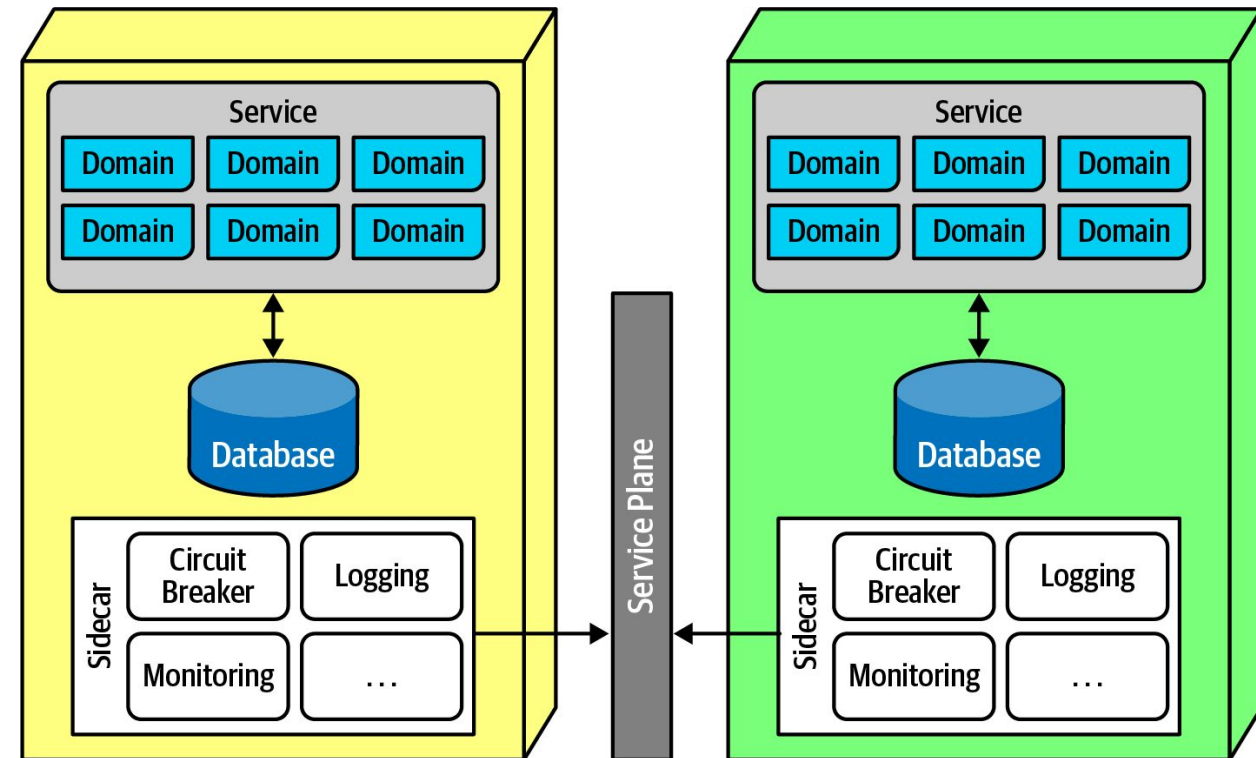


# Boundaries

- Micro is just a label, not a description!
  - In contrast with large services, quite common in 2014!
- How to choose the “right” service granularity?
  - Purpose: domain driven, provides one significant behavior
  - Transactions: not across service boundaries!
  - Choreography: avoid extensive communications among services
- Data isolation is a must, each microservice has its own database
  - There is no single source of truth (old relational DBs)
- The API layer is optional but extremely common
  - Be careful: it is not a mediator or orchestrator!
  - Just a proxy redirecting the request to the “right” microservice

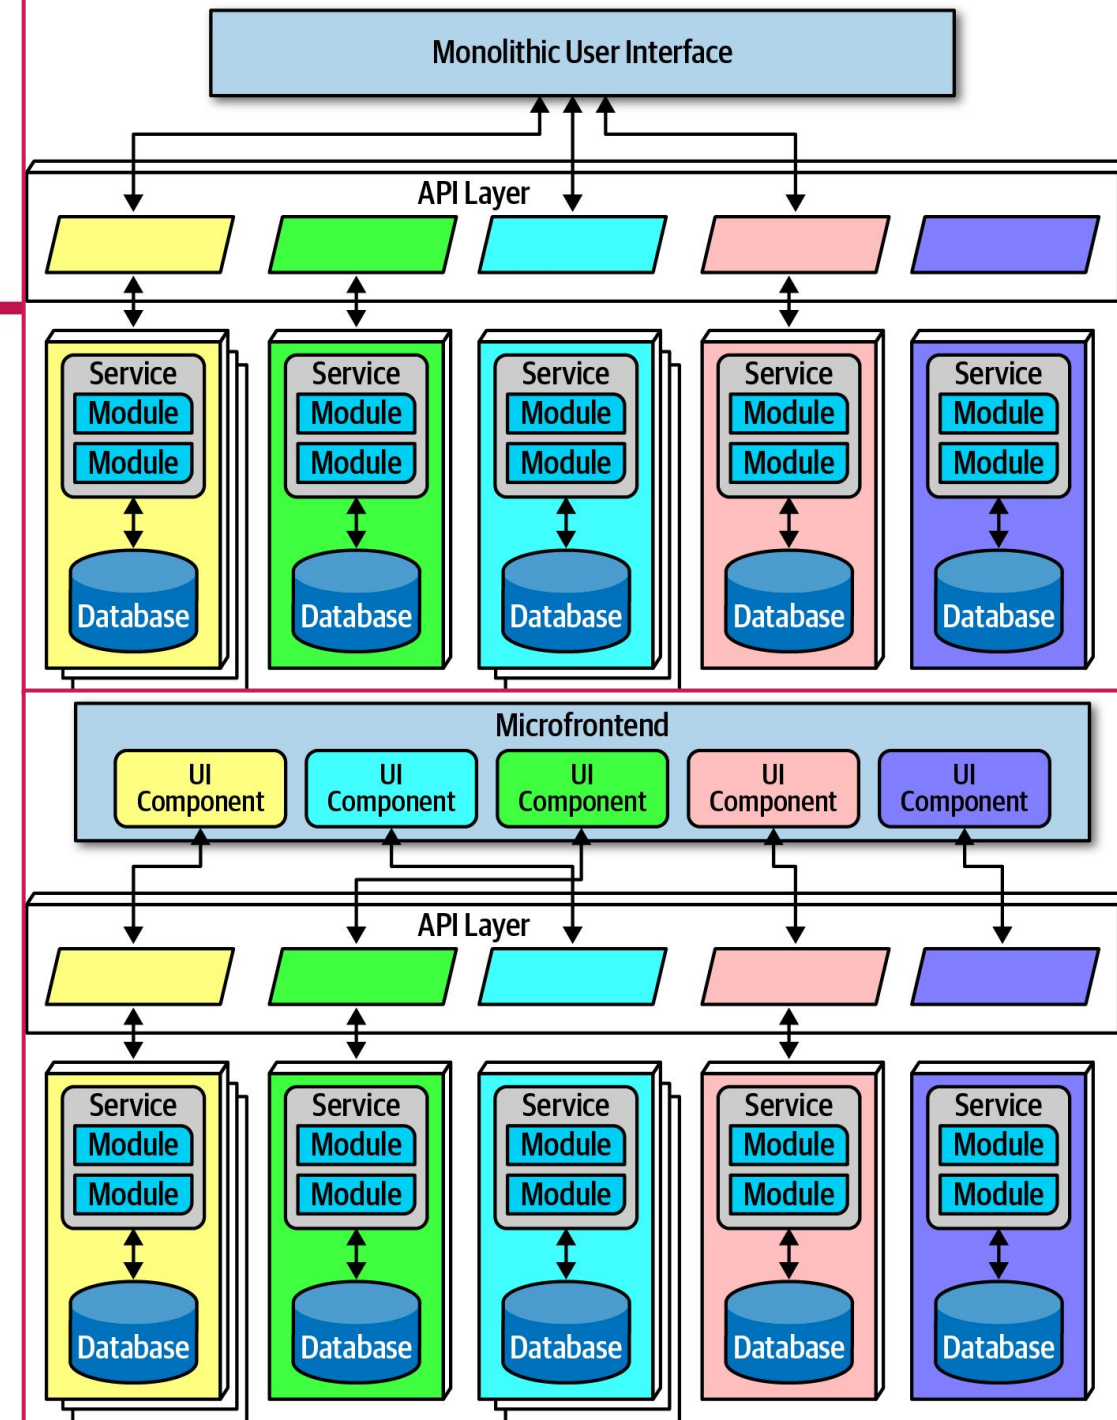
# Operational reuse

- Still shared operational components
  - Logging, monitoring, etc etc...
  - This is unavoidable!
- Some services must be uniform
  - Otherwise the overall project cannot be managed (monitor, deploy, upgrade...)
- A sidecar component is added
  - Allow to build up a service mesh
- Each service is a node in the mesh
  - Console that allows to monitor services
- Service discovery is key for elasticity
  - No direct request, go through API layer
  - Monitor requests, spin up new instances



# Frontends

- UI communicates with many services
  - As service-based architecture
- This is called monolithic frontend
  - Single UI interfaced with the API
- Originally, UIs should be bounded
  - Following domain-driven principles
  - Isolate UIs together with back-end
- Microfrontends are a hot trend
  - Unify entire domain in a single team



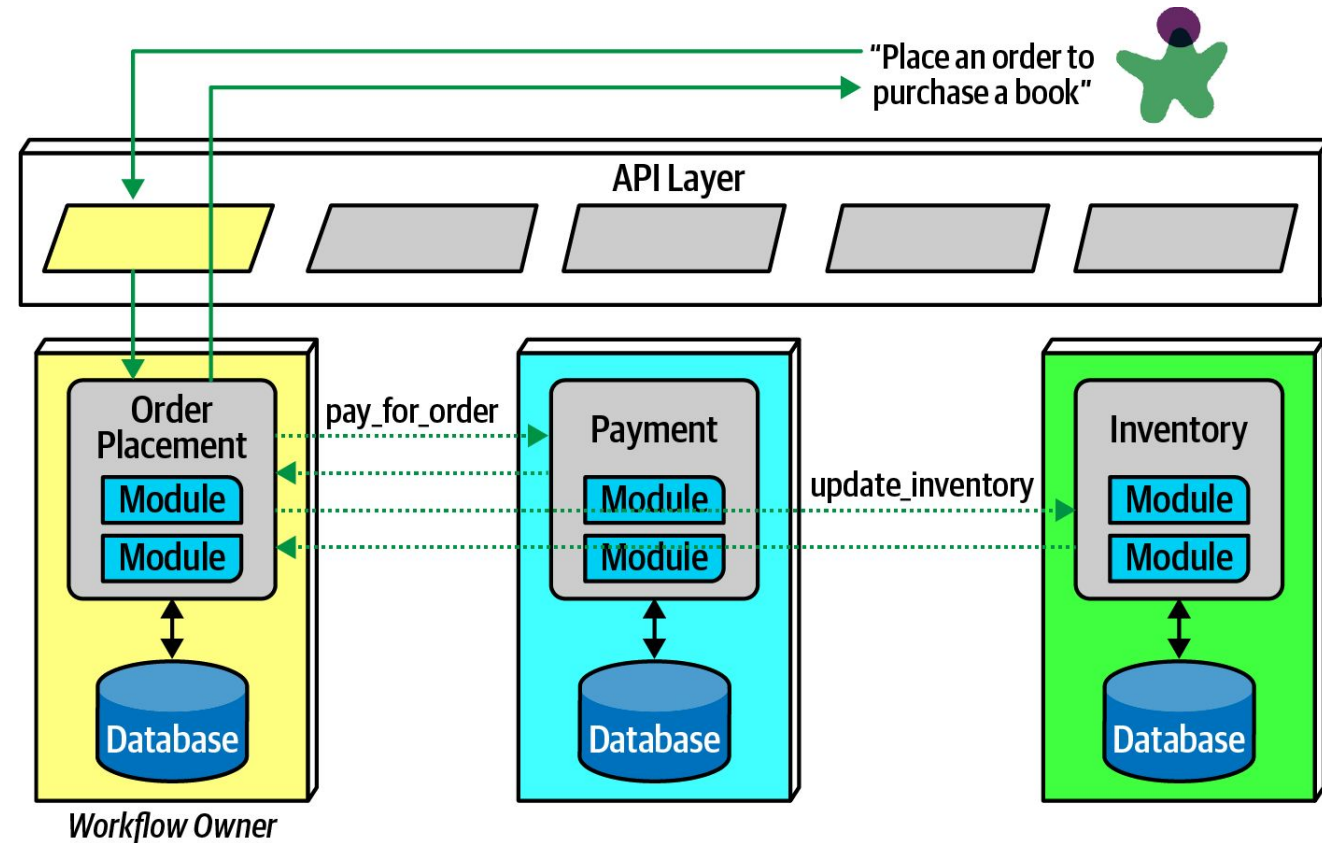


# Synchronous or asynchronous?

- But... what type of communications we have in microservices?
  - In particular, should it be synchronous or asynchronous?
- Each microservice is independent also in this regard
- Protocol-aware heterogeneous interoperability:
  - Protocol-aware: no centralized hub, each service should know how to use other services (e.g., REST APIs, queues, ...)
  - Heterogeneous: each service might rely on different technologies
  - Interoperability: services call one another
- Synchronous communications usually through REST APIs
- Asynchronous communication usually through events/messages

# Choreography

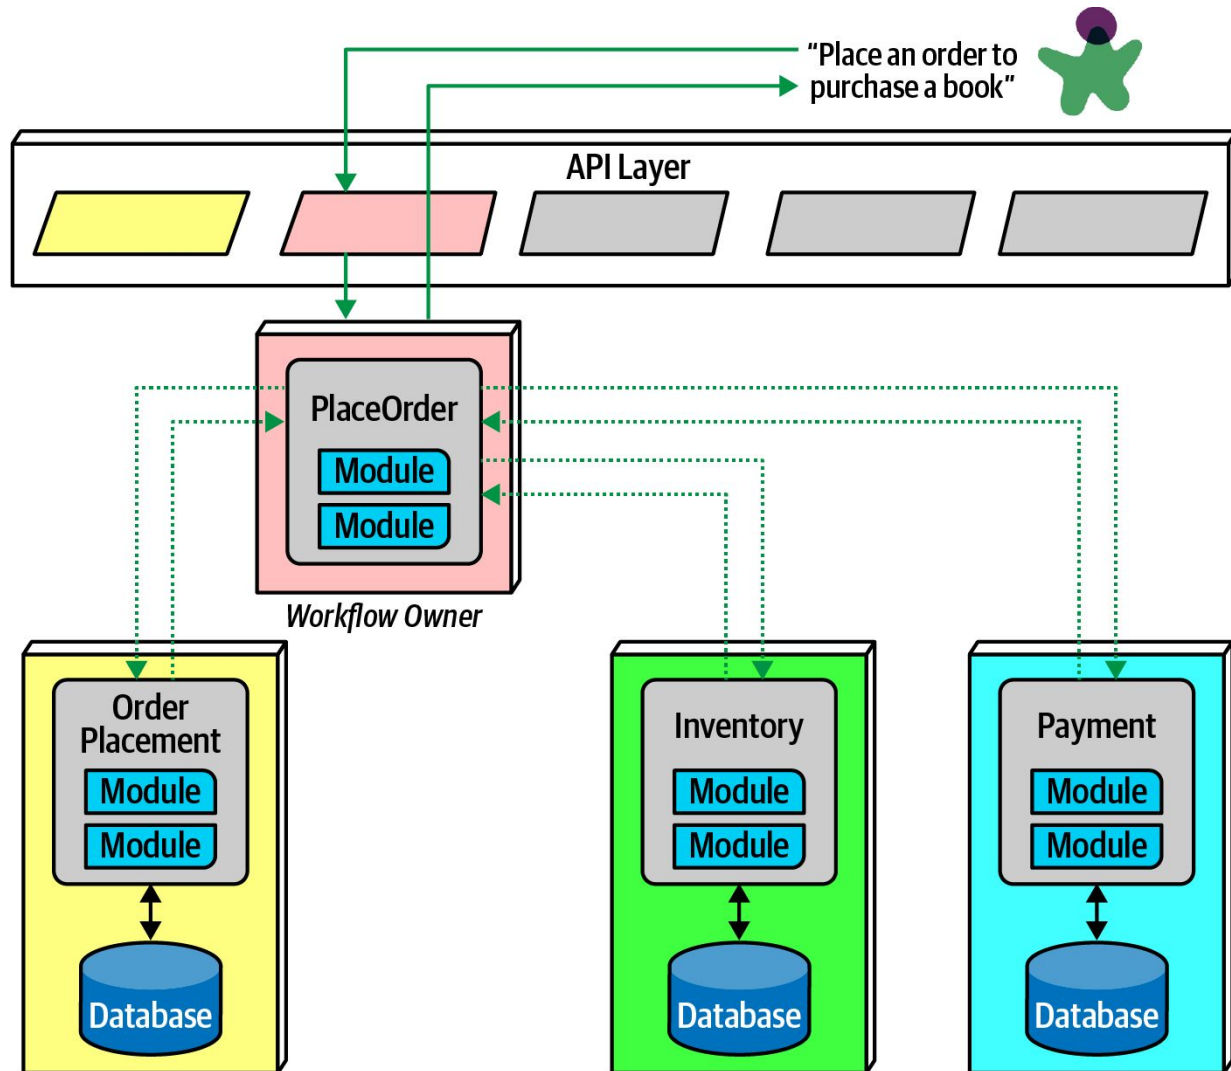
- No central coordinator
  - Broker event-driven architecture
  - Bounded context philosophy
- Each service calls other services as needed
  - This might go pretty deep
- Natural solution aligned with microservice philosophy
- But what happens if failure happens in the middle?
  - See saga





# Orchestration

- We cannot have a global mediator
- Create a service as local mediator
  - Invokes several microservices
  - Coordinate their invocation
  - Aggregate and return the results
- Create coupling between services
  - Mediator depends on all others
  - Sometimes unavoidable
- Single services do not need to coordinate with other services



# Your wooclap poll will be displayed here



Install the **Chrome** or  
**Firefox extension**



Make sure you are in  
**presentation mode**

**wooclap**





Ca' Foscari  
University  
of Venice

---



# Microservices architecture style

## Textbook solution

<b>Deployability</b>	★★★★	<b>Performance</b>	★★
<b>Elasticity</b>	★★★★★	<b>Reliability</b>	★★★★★
<b>Evolutionary</b>	★★★★★	<b>Scalability</b>	★★★★★
<b>Fault tolerance</b>	★★★★★	<b>Simplicity</b>	★
<b>Modularity</b>	★★★★★	<b>Testability</b>	★★★★★
<b>Overall cost</b>	★		



# References

---

- Textbook, Part II
  - Service-based architecture: chapter 13
  - Event-driven architecture: chapter 14
  - Microservices: chapter 17